

Fondamentaux théoriques du Machine Learning (Cours Epita 2022)

NICOLAS LE HIR

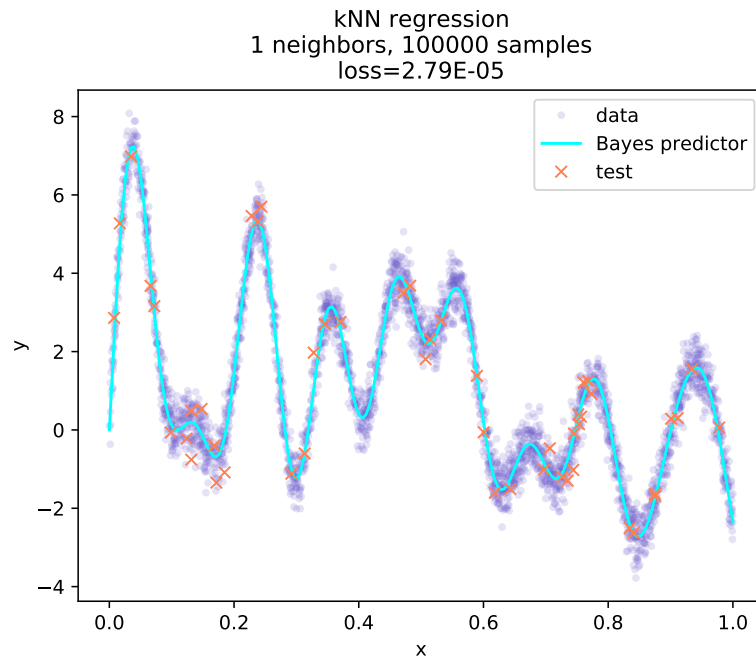


TABLE DES MATIÈRES

1	General mathematical results	3
1.1	Linear algebra	3
1.1.1	Matricial writing of inner products	3
1.2	Statistics and Probability theory	4
1.2.1	Expected value, variance	4
1.2.2	Sample mean, sample variance, sample covariance	6
1.2.3	Markov and Chebychev inequality	7
1.2.4	Concentration inequalities	8
1.3	Differential calculus	8
1.3.1	Differentiable functions	8
1.3.2	Jacobian matrix, gradient and Hessian	9
1.3.3	Lipshitz continuity	10
1.3.4	Smoothness	10
1.4	Algorithmic complexities	11
2	Optimization	11
2.1	Definitions	11
2.2	Existence result	11

2.3	Convex analysis	11
2.3.1	Definitions	12
2.3.2	Differential formulation of convexity	12
2.3.3	Minima of convex functions	13
2.3.4	Existence of a global minima for convex functions	13
2.3.5	Local minimum of a two-times differentiable function	14
2.3.6	Expected values and convex functions	14
2.4	Gradient Descent (GD)	14
2.4.1	Context	14
2.4.2	Minimisation of a strongly convex function	14
2.4.3	Minimization of a convex function	15
3	Supervised Learning	15
3.1	General definitions and bounds	15
3.1.1	Setup	15
3.1.2	Definitions	16
3.1.3	Bayes predictor and conditional expectation	17
3.1.4	Bias-Variance decomposition	18
3.1.5	Expected value of empirical risk	18
3.1.6	Deterministic bound on the estimation error	19
3.2	Probably Approximately Correct Learning : PAC bounds	19
3.2.1	Definition	19
3.2.2	Bound in probability on the fluctuation error	19
3.2.3	Probabilistic bound on the generalization error for a decaying approximation error	20
3.2.4	Approximation error bound	20
3.2.5	Curse of dimensionality	21
3.3	Ordinary Least Squares (OLS)	21
3.3.1	Context	21
3.3.2	OLS estimator	21
3.3.3	Statistical analysis of OLS	22
3.4	Ridge regression	24
3.4.1	Context	24
3.4.2	Ridge regression estimator	24
3.4.3	Statistical analysis of Ridge regression	25
3.4.4	Choice of λ	26
3.5	Lasso	27
3.5.1	Context	27
3.5.2	LASSO estimator	27
3.5.3	Numerical resolutions	27
3.6	Logistic regression	27
3.6.1	Context	27
3.6.2	Logistic regression estimator	28
3.7	Stochastic Gradient Descent (SGD)	28
3.7.1	Context	28
3.7.2	SGD updates	29
3.7.3	SGD as an estimation of GD	29
3.7.4	Convergence for a convex objective	29
3.7.5	Convergence for a strongly convex objective	30
3.7.6	Convergence for a strongly convex L-smooth objective	30
3.7.7	Computational complexity of GD and SGD	30
3.8	Kernels	31
3.8.1	Context	31
3.8.2	Representer theorem	31
3.8.3	Consequence	32
4	Unsupervised learning	33
4.1	Definitions	33

4.1.1	Setup	33
4.2	K-means clustering	34
4.2.1	Context	34
4.2.2	Distorsion	34
4.3	Principal Component Analysis (PCA)	34
4.3.1	Context	34
4.3.2	Optimal projection	34
4.3.3	Algorithm	34
4.4	Density etimation	35
4.4.1	Context	35
4.4.2	Kernel density estimation (KDE) and Convolution	35
5	Probabilistic modeling	35
5.1	Context	35
5.2	Maximum likelihood estimation	36
5.2.1	Link with empirical risk minimization	36
5.2.2	Link with Kullback-Leibler divergence	36
5.2.3	Conditional modeling	36

PRESENTATION

This document gathers important results and tools that will be used during the FTML course at Epita. It is not intended as an exhaustive course on the topic, as some results are stated without proof. It should rather be seen as a compact resource of facts, definitions and results.

1 GENERAL MATHEMATICAL RESULTS

1.1 Linear algebra

1.1.1 Matricial writing of inner products

In machine learning, optimization or statistics we often write the inner product of two vectors of \mathbb{R}^d as a product of matrices. If $x \in \mathbb{R}^d$ writes :

$$x = \begin{pmatrix} x_1 \\ \dots \\ x_i \\ \dots \\ x_d \end{pmatrix}$$

And (with T denoting the transposition),

$$y^T = (y_1, \dots, y_j, \dots, y_d)$$

Then we have that

$$\langle x, y \rangle = y^T x = x^T y$$

For instance, with this property we can show in a compact way that if $y \in \mathbb{R}^n$, $A \in \mathbb{R}^{n \times d}$, and $\theta \in \mathbb{R}^d$, then

$$\langle y, A\theta \rangle = \langle A^T y, \theta \rangle$$

Indeed,

$$\begin{aligned}\langle y, A\theta \rangle &= (A\theta)^T y \\ &= \theta^T A^T y \\ &= \theta^T (A^T y) \\ &= \langle \theta, A^T y \rangle\end{aligned}$$

1.2 Statistics and Probability theory

1.2.1 Expected value, variance

Definition 1. Moments of a distribution

Let X be a real random variable, and $k \in \mathbb{N}^*$. X is said to have a moment of order k if $E(|X|^k) < +\infty$, which means that :

— if X is discrete, with image $X(\Omega) = (x_i)_{i \in \mathbb{N}}$, the series

$$\sum (x_i)^k P(X = x_i)$$

is **absolutely** convergent. The moment is then equal to the sum of that series (without absolute value).

— if X is continuous with density $p(x)$, the integral

$$\int_{-\infty}^{+\infty} x^k f(x) dx$$

is **absolutely** convergent. The moment is then equal to the sum of that series (without absolute value).

Proposition 1. Let $k_1 < k_2$ be integers. Let X be a real random variable. Then if X has a moment of order k_2 , X also has a moment of order k_1 .

Definition 2. Expected value, variance

- If X has a moment of order 1, it is called the **expected value**
- If X has a moment of order 2, then $X - E(X)$ also has a moment of order 2. This moment is called the variance of X .

$$V(X) = E((X - E(X))^2)$$

We often note $\sigma(X) = \sqrt{\text{Var}(X)}$.

Proposition 2. Let a and b be real numbers, and X a random variable that admits a moment of order 2. Then

$$\text{Var}(aX + b) = a^2 \text{Var}(X)$$

Proposition 3. Let X and Y be real two random variables that admit a moment of order 2. Then the variable XY has a moment of order 1, and we then define the covariance of X and Y as

$$\text{Cov}(X, Y) = E((X - E(X))(Y - E(Y)))$$

Lemme 1. Let $X, Y, Z \in \mathbb{R}$ be real random variables with a moment of order 2. We have :

$$\text{Cov}(X + Y, Z) = \text{Cov}(X, Z) + \text{Cov}(Y, Z)$$

$$\text{Var}(X + Y) = \text{Var}(X) + \text{Var}(Y) + 2\text{Cov}(X, Y)$$

$$|\text{Cov}(X, Y)| \leq \sigma(X)\sigma(Y)$$

Proposition 4. Let (X_1, \dots, X_n) be n mutually independent real random variables. Then if they all admit a moment of order 1, then the product $X_1 X_2 \dots X_n$ also does admit a moment of order 1 and

$$E(X_1 X_2 \dots X_n) = \prod_{i=1}^n E(X_i)$$

If they also admit moments of order 2, then

$$\text{Var}\left(\sum_{i=1}^n X_i\right) = \sum_{i=1}^n \text{Var}(X_i)$$

Definition 3. Expected value for vectors, variance matrix

From now on, if we write $E(X)$ or $\text{Var}(X)$, we implicitly assume that the quantities are correctly defined. Let $X \in \mathbb{R}^d$ be a random vector.

$$X = \begin{pmatrix} X_1 \\ \dots \\ X_i \\ \dots \\ X_d \end{pmatrix}$$

The **expected value** of the vector writes

$$E(X) = \begin{pmatrix} E[X_1] \\ \dots \\ E[X_i] \\ \dots \\ E[X_d] \end{pmatrix}$$

Note that we could similarly define the expected value of a matrix.

The **variance matrix** (or **covariance matrix**, **variance-covariance**, **dispersion matrix**) $\text{Var}(X)$ is defined as

$$[\text{Var}(X)]_{ij} = \text{Cov}(X_i, X_j)$$

Remark. It is useful to note that

$$\text{Var}(X) = E\left[(X - E[X])(X - E[X])^T\right] \quad (1)$$

Lemme 2. Let $X \in \mathbb{R}^{d,p}$.

If $A \in \mathbb{R}^{n,d}$, $AX \in \mathbb{R}^{n,p}$, and we have

$$E(AX) = AE(X)$$

If $B \in \mathbb{R}^{p,m}$, $XB \in \mathbb{R}^{d,m}$ and we have

$$E(XB) = E(X)B$$

Lemme 3.

$$\text{Var}(AX) = A\text{Var}(X)A^T \in \mathbb{R}^{n,n}$$

Definition 4. Independent and identically distributed variables

We say that the random variables $(X_n)_{n \in \mathbb{N}}$ are independent and identically distributed if $\forall i \in \mathbb{N}$, the law of X_i is the same as the law of X_1 and they are mutually independent. This is noted **i.i.d variables**.

Remark. In machine learning, when studying a sequence of variables, each one representing for instance a sample drawn from a distribution, we most often assume that the resulting sequence is i.i.d.

1.2.2 Sample mean, sample variance, sample covariance

We observe a dataset $\{x_i \in \mathbb{R}^d, i \in [1, n]\}$ of n i.i.d samples. We assume that the underlying distribution has a moment of order 1, noted $m \in \mathbb{R}^d$.

Definition 5. Sample mean vector

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i \in \mathbb{R}^d$$

Remark. The sample mean vector is an **unbiased estimator** of the expected value.

$$E(\bar{x}) = m$$

Definition 6. Unbiased sample variance

If $d = 1$, we define the **unbiased sample variance** by

$$s^2 = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2$$

Remark. We assume that the underlying distribution has a moment of order 2, noted σ^2 . Then the unbiased sample variance is an **unbiased estimator** of the variance.

$$E(s^2) = \sigma^2$$

If we replace the $\frac{1}{n-1}$ factor by $\frac{1}{n}$, we have a biased estimator. The use of $\frac{1}{n-1}$ is called **Bessel's correction**.

Definition 7. Sample covariance matrix

We go back to the case of $\{x_i \in \mathbb{R}^d, i \in [1, n]\}$. Each sample x_i writes

$$x_i = \begin{pmatrix} x_{i1} \\ \dots \\ x_{ij} \\ \dots \\ x_{id} \end{pmatrix}$$

The sample covariance matrix is a $d \times d$ matrix S^2 such that for all (p, q)

$$S_{pq} = \frac{1}{n-1} \sum_{i=1}^n (x_{ip} - \bar{x}_p)(x_{iq} - \bar{x}_q)$$

This can also be written

$$S = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})(x_i - \bar{x})^T$$

Remark. The sample covariance matrix can also be written in a different way. We consider the **sample matrix** $X \in \mathbb{R}^{n,d}$.

$$X = \begin{pmatrix} x_1^T \\ \dots \\ x_i^T \\ \dots \\ x_n^T \end{pmatrix}$$

We can also consider the **centered sample matrix** defined by

$$X_{centered} = \begin{pmatrix} x_1^T - \bar{x}^T \\ \dots \\ x_i^T - \bar{x}^T \\ \dots \\ x_n^T - \bar{x}^T \end{pmatrix}$$

Then, the sample covariance matrix writes :

$$S^2 = \frac{1}{n-1} X_{centered}^T X_{centered}$$

1.2.3 Markov and Chebychev inequality

Proposition 5. *Markov inequality*

Let X be a real non-negative random variable (variable aléatoire réelle positive), such that $E(|X|) < +\infty$. Let $a > 0$. Then

$$P(X \geq a) \leq \frac{E(X)}{a}$$

Proposition 6. *Chebyshev inequality* Let X be a real random variable, such that $E(|X|^2) < +\infty$. Let $a > 0$. Then

$$P(|X - E[X]| > a) \leq \frac{\text{Var}(X)}{a^2}$$

Theorem 1. *Weak law of large numbers*

Let $(X_n)_{n \in \mathbb{N}}$ be a sequence of i.i.d. variables that have a moment of order 2. We note m their expected value. Then

$$\forall \epsilon > 0, \lim_{n \rightarrow +\infty} P\left(\left|\frac{1}{n} \sum_{i=1}^n X_i - m\right| \geq \epsilon\right) = 0$$

We say that we have **convergence in probability**.

Démonstration. Let $S_n = \frac{1}{n} \sum_{i=1}^n X_i$. By linearity of the expected value, $E(S_n) = m$. We note $\sigma^2 = \text{Var}(X_1)$. By independence of the sequence $(X_n)_{n \in \mathbb{N}}$, we have that $\text{Var}(S_n) = \sum_{i=1}^n \text{Var}(\frac{1}{n} X_i)$, which is equal to $\frac{1}{n^2} \sum_{i=1}^n \text{Var}(X_i) = \frac{\sigma^2}{n}$ with 2.

Using the Chebychev inequality,

$$\begin{aligned} P(|S_n - m| > \epsilon) &\leq \frac{\text{Var}(S_n)}{\epsilon^2} \\ &= \frac{\sigma^2}{n\epsilon^2} \end{aligned}$$

□

Remark. Keeping the same hypotheses on the sequence of variables, we can make some important observations.

- $\lim_{n \rightarrow +\infty} \text{Var}(S_n) = 0$.
- More precisely, we have a probabilistic bound on the magnitude of the difference between the average S_n and m .

$$\sqrt{\text{Var}(S_n - m)} = \frac{\sigma}{\sqrt{n}}$$

The following theorem gives a result that applies to a different type of convergence of random variables (convergence in distribution). As convergence in probability implies convergence in distribution, we say that convergence in distribution is a **weaker** convergence.

Theorem 2. *Central limit theorem* Let $(X_n)_{n \in \mathbb{N}}$ be a sequence of i.i.d. variables that have a moment of order 2. We note m their expected value, and σ^2 their variance. If we consider the following random variable :

$$T_n = \frac{\sqrt{n}}{\sigma} (S_n - m) \tag{2}$$

then T_n converges in distribution to $\mathcal{N}(0, 1)$.

1.2.4 Concentration inequalities

Concentration inequalities are useful in order to have **non-asymptotic** bounds on the deviation between the average of random variables and their mean. Non-asymptotic means that we have an inequality for all n , as opposed to the central limit theorem that is asymptotic and holds for $n \rightarrow +\infty$.

Theorem 3. Bernstein inequality Let u_1, \dots, u_n be i.i.d. random variables, such that $E(u) = 0$ and $|u| \leq B$ almost surely, with $B > 0$. Let $\sigma^2 = E(u^2)$ and $t > 0$:

$$P\left(\frac{1}{n} \sum_{i=1}^n u_i > t\right) \leq e^{-\frac{t^2 n/2}{\sigma^2 + bt}}$$

Corollary 1. Under the same assumptions :

$$P\left(\left|\frac{1}{n} \sum_{i=1}^n u_i\right| > t\right) \leq 2e^{-\frac{t^2 n/2}{\sigma^2 + bt}}$$

Theorem 4. Hoeffding's inequality

Let $(X_i)_{1 \leq i \leq n}$ be n i.i.d real random variables such that $\forall i \in [1, n]$, $X_i \in [a, b]$ and $E(X_i) = \mu \in \mathbb{R}$. Let $\bar{\mu} = \frac{1}{n} \sum_{i=1}^n X_i$.

Then $\forall \epsilon > 0$,

$$P\left(|\bar{\mu} - \mu| \geq \epsilon\right) \leq 2 \exp\left(-\frac{2n\epsilon^2}{(b-a)^2}\right)$$

1.3 Differential calculus

1.3.1 Differentiable functions

Definition 8. Differentiable function

Let V and W be real Hilbert spaces (complete vector space with an inner product). Let $f : V \rightarrow W$. We say that f is differentiable in $x \in V$ if there exists a continuous linear application $L_x : V \rightarrow W$ such that

$$f(x+h) = f(x) + L_x(h) + o(h)$$

$$\text{with } \lim_{h \rightarrow 0} \frac{|o(h)|}{\|h\|} = 0.$$

Remark. We consider the case (that we will often encounter) where $W = \mathbb{R}$. Then for a given $x \in B$, L_x is a linear form, and since V is a Hilbert space, with Riesz representation theorem we know that, there exists a unique vector $p_x \in V$ such that $\forall h \in V$, $L_x(h) = \langle p, h \rangle$. p is sometimes noted $f'(x)$, $\nabla_x f$ or $\nabla f(x)$. We note that we do not assume here that V is of finite dimension.

Definition 9. Two times differentiable function

We keep the same hypotheses as in 8 and assume that $W = \mathbb{R}$. If the application $x \mapsto \nabla_x f$ is differentiable in x , then we say that f is two times differentiable in x . In that case we note $f''(x)$ the second-order derivative, that satisfies :

$$\nabla_{x+h} f = \nabla_x f + f''(x)(h) + o(h)$$

Lemme 4. We keep the same hypotheses. Given x and h , $f''(x)(h)$ is an element of V , that can also be identified to an element of its dual space V^* . With the notation $f''(x)(h, h') = f''(x)(h)(h')$, we can show that

$$f(x+h) = f(x) + \nabla_x f(h) + \frac{1}{2} f''(x)(h, h) + o(\|h\|^2)$$

1.3.2 *Jacobian matrix, gradient and Hessian*

In this paragraph, we consider the case when $V = \mathbb{R}^d$, which is the most frequent situation encountered in machine learning.

Jacobian matrix : If $f : \mathbb{R}^d \rightarrow \mathbb{R}^p$ is differentiable on \mathbb{R}^d we note L_x^f the differential in x . It is an application from \mathbb{R}^d to \mathbb{R}^p . We can then equivalently consider its matrix, that can be noted also $L_x^f \in \mathbb{R}^{p,d}$.

$$f(x+h) = f(x) + L_x^f h + o(h)$$

L_x^f is called the **Jacobian matrix** of f in x .

Link with the gradient : we note that if the function f has real values ($p = 1$), then

$$\nabla_x f = (L_x^f)^T \in \mathbb{R}^{d,1}$$

Jacobian of a product.

If we now consider $g : \mathbb{R}^p \rightarrow \mathbb{R}^q$ another differentiable application, and note its matrix / differential in $f(x)$ as $L_{f(x)}^g \in \mathbb{R}^{q,p}$, we have :

$$\begin{aligned} (g \circ f)(x+h) &= g(f(x+h)) \\ &= g(f(x) + L_x^f h + o(h)) \\ &= g(f(x) + L_x^f h) + L_{f(x)}^g o(h) + o(o(h)) \end{aligned}$$

We can show that $L_{f(x)}^g o(h) + o(o(h)) = o(h)$. Hence,

$$\begin{aligned} (g \circ f)(x+h) &= g(f(x) + L_x^f h) + o(h) \\ &= g(f(x)) + L_{f(x)}^g L_x^f h + o(L_x^f h) + o(h) \end{aligned}$$

We can also show that $o(L_x^f h) = o(h)$, for instance using the operator norm of L_x^f .

Finally, the Jacobian matrix $L_x^{g \circ f}$ of $g \circ f : \mathbb{R}^d \rightarrow \mathbb{R}^q$ in x is the product of the jacobians.

$$L_x^{g \circ f} = L_{f(x)}^g L_x^f \in \mathbb{R}^{q,d} \quad (3)$$

Gradient of a product : With the same functions g and f and if $q = 1$, we deduce from 3 that

$$\begin{aligned} \nabla_x (g \circ f) &= (L_x^{g \circ f})^T \\ &= (L_{f(x)}^g L_x^f)^T \\ &= (L_x^f)^T (L_{f(x)}^g)^T \\ &= (L_x^f)^T \nabla_{f(x)} g \in \mathbb{R}^{d,1} \end{aligned}$$

Hessian : if $f : \mathbb{R}^d \rightarrow \mathbb{R}$ is two times differentiable in x , then the application $\nabla f : \mathbb{R}^d \rightarrow \mathbb{R}^d$, $x \mapsto \nabla_x f$ has a matrix $H_x^f \in \mathbb{R}^{d,d}$, called the Hessian.

$$\nabla_{x+h} f = \nabla_x f + H_x^f h + o(h)$$

Then, the development of f around x can be written

$$f(x+h) = f(x) + L_x^f h + \frac{1}{2} h^T (H_x^f) h + o(\|h\|^2)$$

Examples

Quadratic function Let $A \in \mathbb{R}^{d,d}$ be a symmetric real matrix. If $f(x) = \frac{1}{2} x^T A x - b^T x$, then

$$\text{— } \nabla_x f = A x - b$$

$$— H_x^f = A.$$

Least squares Let $X \in \mathbb{R}^{n,d}$, $\theta \in \mathbb{R}^d$, $y \in \mathbb{R}^n$. If $f(\theta) = \frac{1}{2} \|X\theta - y\|_2^2$, then

$$— \nabla_\theta f = X^T(X\theta - y)$$

$$— H_\theta^f = X^T X.$$

Explicit formulation of the gradient, the Jacobian, the Hessian. We assume $f : \mathbb{R}^d \rightarrow \mathbb{R}^p$ is differentiable.

$$\forall x \in \mathbb{R}^d, f(x) = \begin{pmatrix} f_1(x) \\ \dots \\ f_i(x) \\ \dots \\ f_d(x) \end{pmatrix}$$

$$L_x^f = \begin{pmatrix} \frac{\partial f_1}{\partial x_1}(x) & \frac{\partial f_1}{\partial x_2}(x) & \dots & \frac{\partial f_1}{\partial x_d}(x) \\ \dots & \dots & \dots & \dots \\ \dots & \dots & \dots & \dots \\ \frac{\partial f_p}{\partial x_1}(x) & \frac{\partial f_p}{\partial x_2}(x) & \dots & \frac{\partial f_p}{\partial x_d}(x) \end{pmatrix}$$

If f has real values ($p = 1$), then

$$\nabla_x f = \begin{pmatrix} \frac{\partial f}{\partial x_1}(x) \\ \dots \\ \frac{\partial f}{\partial x_i}(x) \\ \dots \\ \frac{\partial f}{\partial x_d}(x) \end{pmatrix}$$

and if we keep $p = 1$ and assume that f is two times differentiable, then the Hessian reads :

$$H_x^f = \begin{pmatrix} \frac{\partial^2 f}{\partial x_1^2}(x) & \frac{\partial^2 f}{\partial x_2 \partial x_1}(x) & \dots & \frac{\partial^2 f}{\partial x_d \partial x_1}(x) \\ \dots & \dots & \dots & \dots \\ \dots & \dots & \dots & \dots \\ \frac{\partial^2 f}{\partial x_1 \partial x_d}(x) & \frac{\partial^2 f}{\partial x_2 \partial x_d}(x) & \dots & \frac{\partial^2 f}{\partial x_d^2}(x) \end{pmatrix}$$

Remark. By Swartz theorem, if f is C^2 , then the Hessian is symmetrical.

1.3.3 *Lipshitz continuity*

Definition 10. L-Lipschitz continuous function

Let f be a differentiable function and $L > 0$. We say that f is L-Lipschitz continuous if $\forall x, y \in \mathbb{R}^d$,

$$\|f(x) - f(y)\| \leq L\|x - y\|$$

Definition 11. L-Lipschitz continuous gradients

Let f be a differentiable function and $L > 0$. We say that f has L-Lipschitz continuous gradients if $\forall x, y \in \mathbb{R}^d$,

$$\|\nabla_x f - \nabla_y f\| \leq L\|x - y\|$$

1.3.4 *Smoothness*

Definition 12. Smoothness

A differentiable function f with real values is said L-smooth if and only if

$$\forall x, y \in \mathbb{R}^d, |f(y) - f(x) - \nabla_x f(y - x)| \leq \frac{L}{2} \|y - x\|^2$$

Lemme 5. f is L-smooth if and only if it has L-Lipshitz continuous gradients.

Lemme 6. If f is twice differentiable, f is L -smooth if and only if

$$\forall x \in \mathbb{R}^d, -LI_d \leq H_x f \leq LI_d$$

Meaning that all eigenvalues of $H_x f$ have a module that is smaller than L . Here, \geq denotes the **Loewner order**, which is a partial order on symmetric matrices. https://en.wikipedia.org/wiki/Loewner_order.

1.4 Algorithmic complexities

In machine learning, there is a focus on the efficiency of algorithm, as large-scale problems are frequent and some algorithmic complexities are prohibitive. We give some reminders on orders of magnitudes of time-complexity of some operations, ignoring space complexity in this section.

- If $A \in \mathbb{R}^{d,d}$ is invertible, the cost of inverting it is $\mathcal{O}(d^3)$.
- If $A \in \mathbb{R}^{d,d}$ is symmetric, the cost of computing its eigenvalues and eigenvectors $\mathcal{O}(d^3)$.
- If $X \in \mathbb{R}^{n,d}$ and $\theta \in \mathbb{R}^d$, the cost of computing $X\theta \in \mathbb{R}^n$ is $\mathcal{O}(nd)$.

Some optimisations of these operations exist, see :

- QR decomposition for matrix inversion (https://en.wikipedia.org/wiki/QR_decomposition).
- Low-rank approximation. https://en.wikipedia.org/wiki/Low-rank_approximation

2 OPTIMIZATION

Reference : [Allaire, 2012,]

2.1 Definitions

Definition 13. Minima Let $f : \mathbb{R}^d \rightarrow \mathbb{R}$ be defined on $K \subset \mathbb{R}^d$.
 $x \in K$ is a local minimum of f on K if and only if

$$\exists \delta > 0, \forall y \in K, \|y - x\| < \delta \Rightarrow f(x) \leq f(y)$$

$x \in K$ is a global minimum of f on K if and only if

$$\forall y \in K, f(x) \leq f(y)$$

Definition 14. Coercive function

$f : V \rightarrow \mathbb{R}$, defined on a vector space, V is **coercive** if

$$\lim_{\|x\| \rightarrow +\infty} f(x) = +\infty$$

2.2 Existence result

Theorem 5. Existence of a global minimum in \mathbb{R}^d

Let K be a closed non-empty subset of \mathbb{R}^d , and $f : \mathbb{R}^d \rightarrow \mathbb{R}$ a continuous coercive function. Then, there exists at least a global minimum of f on K .

2.3 Convex analysis

Convexity is a key property of function, that allows to have theoretical proofs on the existence of solutions to a optimization problem, and on the convergence of optimization algorithms. It is useful in finite dimensional spaces and also in infinite dimensional spaces.

2.3.1 Definitions

Definition 15. Convex function

The function $f : \Omega \rightarrow \mathbb{R}$ with Ω convex is :

— **convex** if $\forall x, y \in \Omega, \alpha \in [0, 1]$,

$$f(\alpha x + (1 - \alpha)y) \leq \alpha f(x) + (1 - \alpha)f(y)$$

— **strictly convex** if $\forall x, y \in \Omega, \alpha \in [0, 1]$,

$$f(\alpha x + (1 - \alpha)y) < \alpha f(x) + (1 - \alpha)f(y)$$

— **μ -strongly convex** if $\forall x, y \in \Omega, \alpha \in [0, 1]$,

$$f(\alpha x + (1 - \alpha)y) \leq \alpha f(x) + (1 - \alpha)f(y) - \frac{\mu}{2}\alpha(1 - \alpha)\|x - y\|^2$$

Lemme 7. Let f be convex on convex set Ω , $f : \Omega \rightarrow \mathbb{R}$. Let $c \in \mathbb{R}$. Then

$$\Omega_c = \{x \in \Omega, f(x) \leq c\}$$

is a convex set.

Proposition 7. f is μ convex if and only if $x \mapsto f(x) - \frac{\mu}{2}\|x\|^2$ is convex.

Examples :

- All norms are convex.
- $x \mapsto \log(1 + e^{-x})$ is convex on \mathbb{R}
- $x \mapsto \theta^T x$ is convex on \mathbb{R}^d with $\theta \in \mathbb{R}^d$ (linear form)
- if Q is a symmetric semidefinite positive matrix (matrice positive), then the positive quadratic form $x \mapsto x^T Q x$ is convex.
- if Q is a symmetric definite positive matrix (matrice définie positive) with smallest eigenvalue $\lambda_{\min} > 0$, then the positive definite quadratic form $x \mapsto x^T Q x$ is $2\lambda_{\min}$ -strongly convex.
- if $X \in \mathbb{R}^{n \times d}$, and $Y \in \mathbb{R}^n$, $\theta \mapsto \|X\theta - Y\|^2$ is convex on \mathbb{R}^d (Least-squares)
- If f is increasing and convex and g is convex, then $f \circ g$ is convex.
- If f is convex and g is linear, then $f \circ g$ is convex.

Definition 16. Condition number

If f is L -smooth and μ -strongly convex, we define its **condition number** as :

$$\kappa = \frac{L}{\mu}$$

Remark. The condition number is an important quantity to consider when studying the convergence speed of gradient algorithms.

2.3.2 Differential formulation of convexity

Proposition 8. Let $f : V \rightarrow \mathbb{R}$ be a differentiable function. The following conditions are equivalent.

- f is convex.
- $\forall x, y \in V, f(y) \geq f(x) + (f'(x)|y - x)$ (f is above its tangent space)
- $\forall x, y \in V, (f'(x) - f'(y)|x - y) \geq 0$ (f' grows)

Remark. Equivalently

$$\forall u, v \in V, f(u) - f(v) \leq f'(u)(v - u)$$

Proposition 9. Let $f : V \rightarrow \mathbb{R}$ be a differentiable function, and $\mu > 0$. The following conditions are equivalent.

- f is μ -convex
- $\forall x, y \in V, f(y) \geq f(x) + (f'(x)|y - x) + \frac{\mu}{2}\|y - x\|^2$
- $\forall x, y \in V, (f'(x) - f'(y)|x - y) \geq \mu\|x - y\|^2$

Lemme 8. *Łojasiewicz inequality*

Let $f : \mathbb{R}^d \rightarrow \mathbb{R}$ be a μ -strongly convex differentiable function with a unique minimiser x^* . Then, we have that for all $x \in \mathbb{R}^d$,

$$\|\nabla_x f\|_2^2 \geq 2\mu(f(x) - f(x^*))$$

Proposition 10. *Convexity of two-times differentiable functions*

We now assume that f is two times differentiable.

- f is convex if and only if

$$\forall x, h \in y, J''(x)(h, h) \geq 0$$
- f is μ -strongly convex if and only if

$$\forall x, h \in y, J''(x)(h, h) \geq \mu\|h\|^2$$

Remark. In the case of $V = \mathbb{R}^d$, the two previous conditions mean respectively that

$$\forall x, h \in y, h^T(H_x f)h \geq 0 \tag{4}$$

and

$$\forall x, h \in y, h^T(H_x f)h \geq \mu\|h\|^2 \tag{5}$$

4 means that $\forall x \in \mathbb{R}^d$, all eigenvalues of $H_x f$ are non-negative, while 5 means that they all are $\geq \mu$.

2.3.3 Minima of convex functions

Proposition 11. — If f is convex, any local minimum is a global minimum. The set of global minimizers is a convex set.

- If f is strictly convex, there exists at most one local minimum (that is thus global).
- If f is convex and C^1 (differentiable, $a \mapsto df_a$ continuous), then x is a minimum (thus global) of f on V if and only if the gradient cancels in x , $\nabla_x f = 0$. V need not be finite-dimensional.

2.3.4 Existence of a global minima for convex functions

Theorem 6. *Existence of a global minimum, strongly convex case*

Let K be a closed non-empty subset of the Hilbert space H , and $f : \mathbb{R}^d \rightarrow \mathbb{R}$ a continuous μ -strongly convex function. Then, there exists a unique a global minimum u of f on K , and

$$\forall y \in K, \|y - x\|^2 \leq \frac{\alpha}{4}[f(y) - f(x)]$$

Theorem 7. *Existence of a global minimum, convex case*

Let K be a closed non-empty subset of the Hilbert space H , and $f : \mathbb{R}^d \rightarrow \mathbb{R}$ a continuous convex coercive function. Then, there exists a global minimum u of f on K .

2.3.5 Local minimum of a two-times differentiable function

Theorem 8. We assume that $K = V$ and f is two-times differentiable in x .

If x is a local minimizer of f , then

- $f'(x) = 0$
- $\forall h \in V, f''(u)(h, h) \geq 0$

Remark. In this theorem, we do not assume that f is convex.

Remark. When $V = \mathbb{R}^d$, the previous conditions translate in the fact that $H_x f$ is positive semi-definite or that $H_y f$ is positive semidefinite in a neighborhood of x .

2.3.6 Expected values and convex functions

Let D be a convex set. If f is convex, then if $p : S \Rightarrow D$ is such that $p(x) \geq 0$ and $\int_S p(x) dx = 1$, then

$$f\left(\int_S p(x) x dx\right) \leq \int_S p(x) f(x) dx$$

Let D be a convex set. If X is a random variable such that $X \in D$ almost surely and $E[X]$ exists, then

$$f(E[X]) \leq E[f(X)]$$

2.4 Gradient Descent (GD)

2.4.1 Context

GD is an important first-order algorithm for optimization. First-order means that it is based on the first-order derivative. For well-conditioned convex problems, it converges exponentially fast. We want to minimize a differentiable function f . The **gradient descent (GD)** update reads :

$$x_{t+1} = x_t - \gamma_t \nabla f(x_t)$$

2.4.2 Minimisation of a strongly convex function

The following lemma shows that if the step size is small enough, then GD is a **descent algorithm**.

Lemma 9. let $L > 0$, $f : \mathbb{R}^d \rightarrow \mathbb{R}$ a convex function with L -Lipshitz continuous gradients (see 1.3.3). Let $x_0 \in \mathbb{R}^d$, $\gamma_t > 0$. Then

$$f(x_t) \leq f(x_{t-1}) - \gamma_t(1 - L\gamma_t) \|\nabla f(x_{t-1})\|^2$$

In particular if $0 < \gamma_t < \frac{1}{L}$, then $\forall t \in \mathbb{N}$ such that x_{t-1} is not a global minimum,

$$f(x_t) < f(x_{t-1})$$

Otherwise $x_t = x_{t-1}$ and $f(x_t) = f(x_{t-1})$.

Theorem 9. Convergence of GD for a strongly convex function

Let $f : \mathbb{R}^d \Rightarrow \mathbb{R}$ be a μ -strongly convex function with L -Lipshitz continuous gradients. Let x^* be the global minimum of f (which we know exists since f is strongly convex), $x_0 \in \mathbb{R}^d$, $T \in \mathbb{N}$.

- With constant step size $\gamma_t = \frac{1}{L}$, we have

$$f(x_t) - f(x^*) \leq \left(1 - \frac{\mu}{L}\right)^t (f(x_0) - f(x^*))$$

— With constant step size $0 < \gamma < \frac{1}{2L}$, we have

$$\|x_t - x^*\|^2 \leq (1 - \gamma\mu)^T \|x_0 - x^*\|^2 \quad (6)$$

Remark. 6 could seem surprising at first sight, since it implies that $(1 - \mu\gamma) \geq 0$ if $0 < \gamma < \frac{1}{2L}$. However, we have

$$\mu\|x - y\|^2 \leq \langle \nabla_x f - \nabla_y f, x - y \rangle \leq L\|x - y\|^2$$

Hence, $\mu \leq L$, so $\mu\gamma < \frac{1}{2}$.

Corollary 2. We note that $(1 - \gamma\mu)^T = (1 - \frac{\gamma\mu^T}{T})^T \leq e^{-\gamma\mu^T}$. Thus,

$$\|x_t - x^*\|^2 \leq e^{-\gamma\mu^T} \|x_0 - x^*\|^2$$

We call this an **exponential convergence**. It is also sometimes called **linear convergence**. We can deduce that if

$$T \geq \frac{1}{\gamma\mu} \log\left(\frac{\|x_0 - x^*\|^2}{\epsilon}\right)$$

then

$$\|x_T - x^*\| \leq \epsilon$$

2.4.3 Minimization of a convex function

Theorem 10. Convergence of GD for a smooth convex function

Let $f : \mathbb{R}^d \rightarrow \mathbb{R}$, with global minimiser x^* . With constant step-size $\gamma_t = \frac{1}{t}$, the iterates x_t of GD satisfy :

$$f(x_t) - f(x^*) \leq \frac{L}{2t} \|x_0 - x^*\|^2$$

3 SUPERVISED LEARNING

3.1 General definitions and bounds

3.1.1 Setup

Given some **observations** $(x_i, y_i) \in \mathcal{X} \times \mathcal{Y}$ of input/output pairs, we would like to predict well the new output $y \in \mathcal{Y}$ that should be associated with a new input $x \in \mathcal{X}$. The training dataset is noted $D_n = \{(x_i, y_i), i \in [1, \dots, n]\}$.

We assume that there exists an **unknown distribution** ρ for the joint variable (X, Y) . In most theoretical frameworks used to study Supervised Learning, each sample (x_i, y_i) is assumed to be drawn independently from this law. Hence, D_n is a **random variable**.

$$(x_i, y_i) \sim \rho, \forall i \in [1, n]$$

A **learning rule** \mathcal{A} is a application that associates a **prediction function**, or **estimator** \tilde{f}_n , to D_n .

$$\mathcal{A} : \begin{cases} \cup_{n \in \mathbb{N}} (\mathcal{X} \times \mathcal{Y})^n \rightarrow \mathcal{Y}^{\mathcal{X}} \\ D_n \mapsto \tilde{f}_n \end{cases}$$

Since D_n is a random variable, and since \tilde{f}_n depends on D_n , \tilde{f}_n is also a random variable.

3.1.2 Definitions

Definition 17. Loss function

A **loss function** l is an application that measures the discrepancy between to elements of a set (for instance of a linear space).

$$l: \begin{cases} \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}_+ \\ (y, y') \mapsto l(y, y') \end{cases}$$

Examples : The most common loss functions are the following :

"0-1" loss for **classification**. $\mathcal{Y} = \{0, 1\}$ or $\mathcal{Y} = \{-1, 1\}$.

$$l(y, z) = 1_{y \neq z} \quad (7)$$

square loss for **regression**. $\mathcal{Y} = \mathbb{R}$.

$$l(y, z) = (y - z)^2 \quad (8)$$

absolute loss for **regression**. $\mathcal{Y} = \mathbb{R}$.

$$l(y, z) = |y - z| \quad (9)$$

Definition 18. Risks

Let l be a loss.

The **risk** (or **statistical risk**, **generalization error**, **test error**) of estimator f writes

$$\mathbb{E}_{(X,Y) \sim \rho} [l(Y, f(X))]$$

The **empirical risk (ER)** of an estimator f writes

$$R_n(f) = \frac{1}{n} \sum_{i=1}^n l(y_i, f(x_i))$$

We emphasize that the risks depends on the loss l .

Remark.

$$R_n(f) = \mathbb{E}_{(X,Y) \sim \rho_n} [l(Y, f(X))]$$

We define the **target function** f^* by

$$f^* \in \arg \min_{f: X \rightarrow Y} R(f)$$

with $f: X \rightarrow Y$ set of measurable functions.

Definition 19. Fundamental problem of Supervised Learning

Estimate f^* given only D_n and l .

Definition 20. Excess risk

The **excess risk** $\mathcal{R}(\tilde{f}_n)$ measures how close \tilde{f}_n (predictor obtained by ERM) is to the best possible f^* , in terms of expected risk (average / expecter) error on new examples.

$$\mathcal{R}(\tilde{f}_n) = R(\tilde{f}_n) - R(f^*)$$

We also define the **excess probability** as

$$P(\mathcal{R}(\tilde{f}_n) - R(f^*) > t), \forall t > 0$$

Remark. Since \tilde{f}_n is a random variable D_n , $R_n(\tilde{f})$ and thus $\mathcal{R}(\tilde{f}_n)$ are also random variables.

Definition 21. Consistency

The algorithm \mathcal{A} is said to be **consistent** if

$$\lim_{n \rightarrow +\infty} \mathbb{E}_{D_n} \mathcal{R}(\tilde{f}_n) = 0$$

We also say that \mathcal{A} is a **learning algorithm**.

Definition 22. Strong consistency

The algorithm \mathcal{A} is said to be **strongly consistent** if with probability 1

$$\lim_{n \rightarrow +\infty} \mathcal{R}(\tilde{f}_n) = 0$$

Definition 23. Learning Rates

The sequence $(e_n)_{n \in \mathbb{N}}$ is a learning rate in expectation if

$$\mathbb{E}_{D_n}[\mathcal{R}(\tilde{f}_n)] \leq e_n, \forall n \in \mathbb{N}$$

Definition 24. Learning rate in probability

Let $\delta \in]0, 1]$. The sequence $(p_{n,\delta})_{n \in \mathbb{N}}$ is a learning rate in probability if

$$\mathbb{P}_{D_n}(\mathcal{R}(\tilde{f}_n) \geq p_{n,\delta}) \leq \delta, \forall n \in \mathbb{N}$$

3.1.3 Bayes predictor and conditional expectation

Under some conditions, we can give an explicit formulation of f^* , the best predictor in $\mathcal{Y}^{\mathcal{X}}$, although we can not compute it without the knowledge of the distribution of (X, Y) . In this paragraph we ignore the measurability issues that were mentioned earlier.

Theorem 11. Bayes predictor

If \mathcal{Y} is compact and \mathfrak{l} is continuous, we have

$$f^*(x) = \arg \min_{y \in \mathcal{Y}} \mathbb{E}[\mathfrak{l}(Y, y) | X = x] \quad (10)$$

almost everywhere. $\mathbb{E}[q(Y) | X = x]$ denotes the conditional expectation of $q(Y)$ given that $X = x$, with $q : Y \mapsto \mathbb{R}$. f^* is also called the **Bayes predictor**.

Definition 25. Bayes risk

The Bayes risk is the risk of the Bayes predictor.

$$R^* = \mathbb{E}_X \left(\inf_{y \in \mathcal{Y}} \mathbb{E}(\mathfrak{l}(Y, y) | X) \right)$$

Remark. — Y is a random variable while the parameter y in 10 is a scalar.

- As we do not know the joint distribution ρ , it is not possible to compute the Bayes predictor from this expression.
- Here, each possible value x taken by X is considered independently, as we do not require any regularity in $f^* : \mathcal{X} \rightarrow \mathcal{Y}$.

Bayes estimator for regression and squared loss

In that case, we can be more precise on the value of the Bayes predictor. We assume that $\mathcal{X} = \mathbb{R}^d$ and $\mathcal{Y} = \mathbb{R}$, and that we use the square loss $\mathfrak{l}(y, z) = (y - z)^2$. Hence, for each $x \in \mathcal{X}$,

$$f^*(x) = \arg \min_{y \in \mathcal{Y}} \mathbb{E}((y - Y)^2 | X = x)$$

However, given y ,

$$\begin{aligned} E((y - Y)^2 | X = x) &= E((y - E(Y|X = x)) + (E(Y|X = x) - Y))^2 | X = x) \\ &= E((y - E(Y|X = x))^2 + 2(y - E(Y|X = x))(E(Y|X = x) - Y) + (E(Y|X = x) - Y)^2 | X = x) \end{aligned}$$

$(y - E(Y|X = x))^2$ is a scalar, so

$$E(y - E(Y|X = x))^2 = (y - E(Y|X = x))^2$$

We also have that

$$E(2(y - E(Y|X = x))(E(Y|X = x) - Y)) = 2(y - E(Y|X = x))E(E(Y|X = x) - Y) = 0$$

As a consequence, the value that minimizes $E((y - Y)^2 | X = x)$ is $E(Y|X = x)$.

3.1.4 Bias-Variance decomposition

There are several possible **risk decompositions** that can be written. All these decompositions aim at making a distinction between different, often antagonist contributions to the excess risk. We remember that as \tilde{f}_n is a random variable, $R(\tilde{f}_n)$ is also a random variable. One possible risk decomposition is the following one.

$$E[R(\tilde{f}_n)] - R^* = \left(E[R(\tilde{f}_n)] - \inf_{f \in F} R(f) \right) + \left(\inf_{f \in F} R(f) - R^* \right)$$

We note that both terms are positive.

— **Approximation error (bias term)** : depends on f^* and F , not on \tilde{f}_n , D_n .

$$\inf_{f \in F} R(f) - R^*$$

It is sometimes also defined as $\inf_{f \in F} R(f)$.

— **Estimation error (variance term, fluctuation error, stochastic error)** : depends on D_n , F , \tilde{f}_n . To bound this term we need assumptions on the distribution P .

$$E(R(\tilde{f}_n)) - \inf_{f \in F} R(f)$$

When the algorithm is **underfitting**, the capacity of the hypothesis space is too small, the functions are not able to approximate f^* well and the variance is small. On the contrary, if F is too large, there likely exists a function $f \in F$ close to f^* , inducing a small bias, but the variance is large (**overfitting**).

Remark.

Definition 26. Best estimator in hypothesis space

It is defined as

$$f_a = \arg \min_{f \in F} R(h)$$

Then the bias-variance decomposition can also be written as

$$E[R(\tilde{f}_n) - R^*] = \left(E(R(\tilde{f}_n)) - R(f_a) \right) + \left(R(f_a) - R^* \right)$$

3.1.5 Expected value of empirical risk

For some **fixed** $h \in H$, $R_n(h)$ is a random variable, as it depends on the dataset D_n . let us show that $R_n(h)$ is an **unbiased estimator** of the real risk R . We assume that the samples (X_i, Y_i) are i.i.d, with the distribution of (X, Y) , noted ρ .

$$\begin{aligned}
E_{D_n \sim \rho}(R_n(h)) &= \frac{1}{n} \sum_{i=1}^n E_{D_n \sim \rho}(l(h(X_i), Y_i)) \\
&= \frac{1}{n} \sum_{i=1}^n E_{(X,Y) \sim \rho}(l(h(X), Y)) \\
&= E_{(X,Y) \sim \rho}(l(h(X), Y)) \\
&= R(h)
\end{aligned}$$

However, this is not true for \tilde{f}_n , as \tilde{f}_n depends on the training dataset D_n . We cannot say that for all $i \in [1, n]$, $E_{D_n \sim \rho}(l(\tilde{f}_n(X_i), Y_i)) = E_{(X,Y) \sim \rho}(l(\tilde{f}_n(X), Y))$.

For instance, let us consider the following situation. X follows a uniform law on $]0, 1]$, and $Y = 3X + \sigma\epsilon$, with ϵ being a standard Gaussian random variable independent from X , hence $\mathcal{X} =]0, 1]$ and $\mathcal{Y} = \mathbb{R}$. We assume that $n = 1$, and $D_1 = \{(X_1, Y_1)\} = \{(\alpha, \beta)\}$. We assume that F , the space of available functions for \tilde{f}_n , is the space of linear functions, of the form $x \mapsto \alpha x$, with $\alpha \in \mathbb{R}$. If our learning rule is empirical risk minimization, with quadratic loss, then the estimator \tilde{f}_n writes $\tilde{f}_n(x) = \frac{\beta}{\alpha}x$, and $l(\tilde{f}_n(\alpha), \beta) = 0$. Hence,

$$E_{D_1 \sim \rho}(l(\tilde{f}_n(X_1), Y_1)) = 0 \quad (11)$$

However, we do not have $E_{(X,Y) \sim \rho}(l(\tilde{f}_n(X), Y)) = 0$.

$$\begin{aligned}
E_{(X,Y)}(l(\tilde{f}_n(X), Y)) &= E_{(X,\epsilon)}(l(\tilde{f}_n(X), 3X + \sigma\epsilon)) \\
&= E_{(X,\epsilon)}\left(\left(\left(\frac{\beta}{\alpha} - 3\right)X - \sigma\epsilon\right)^2\right) \\
&\neq 0
\end{aligned}$$

3.1.6 Deterministic bound on the estimation error

Proposition 12.

$$R(f_a) \leq R(\tilde{f}_n) \leq R(f_a) + 2 \sup_{f \in F} |R(f) - R_n(f)|$$

Remark. This can be written as

$$R(\tilde{f}_n) - R(f_a) \leq 2 \sup_{f \in F} |R(f) - R_n(f)|$$

3.2 Probably Approximately Correct Learning : PAC bounds

3.2.1 Definition

Definition 27. PAC accuracy

We say that \tilde{f}_n is ϵ -accurate with confidence $1 - \delta$ or (ϵ, δ) -PAC if

$$P_{D_n} \left(R(\tilde{f}_n) - \inf_{f \in F} R(f) > \epsilon \right) < \delta$$

Remark. Equivalently :

$$P_{D_n} \left(R(\tilde{f}_n) - \inf_{f \in F} R(f) \leq \epsilon \right) \geq 1 - \delta$$

3.2.2 Bound in probability on the fluctuation error

Theorem 12. Assumptions :

— $\exists B > 0, \forall y, y' \in \mathcal{Y}, l(y, y') \leq B$, Hence, $\forall f \in F, R(f) \leq B$.

— F is finite of cardinal $|F|$.

Let $\delta > 0$. If $n \geq \frac{2|F|}{\delta}$, we have :

$$P\left(R(\tilde{f}_n) - \inf_{f \in F} (R(f)) > \sqrt{\frac{4B^2 \log(\frac{2|F|}{\delta})}{n}}\right) \leq \delta$$

Corollary 3. With probability $1 - \delta$, we have

$$R(\tilde{f}) \leq R(f_a) + 2\sqrt{\frac{\log(|F|) + \log(\frac{2}{\delta})}{2n}}$$

Remark. — We say that the estimation error is $\mathcal{O}(\frac{1}{\sqrt{n}})$.

— This bound does not depend on the law p from which the samples are drawn. Thus, it is a bound that is **uniform** with regards to the law of (X, Y) .

— $\log(|F|)$ can be seen as the entropy of the class H , if the probability distribution is uniform on that class.

— theoretically, H is often an **infinite set**. However, it can be discretized. If H has m parameters quantified on N values, then $|F| = N^m$, hence $\log(H) = m \log(N)$. In order to control the term $\frac{\log(|F|)}{2n}$, then we should have $n \gg m$.

— To have results for infinite sets, the Vapnik-Chervonenkis theory is required.

3.2.3 Probabilistic bound on the generalization error for a decaying approximation error

If we add a hypothesis on the **decay** of the approximation error $R(f_a)$, as a function of $|F|$, we can bound the **generalisation error**.

Proposition 13. Assume $\exists \beta, C > 0$ such that

$$R(f_a) \leq C (\log(|F|))^{-\beta} \quad (12)$$

$\forall \epsilon > 0$, if we have

$$n \geq C^{\frac{1}{\beta}} \epsilon^{-2-\frac{1}{\beta}}$$

then,

$$P(R(\tilde{f}_n) \leq 3\epsilon) \geq 1 - 2e^{-C^{\frac{1}{\beta}} \epsilon^{-\frac{1}{\beta}}}$$

3.2.4 Approximation error bound

Noiseless model : We introduce an *a priori* knowledge in the problem, stating that there exists a (unknown) function f linking the inputs x to the outputs, without ambiguity. Thus, we can write $y = f(x)$.

Classes of functions : We also assume f belongs to a given class of functions F . Usually, F determines the **regularity** of f (differentiability order, Lipschitz, etc.).

$$R(f_a) \leq \max_{f \in F} \min_{h \in H} \|f - h\|_{\infty}$$

Thus, if we have bounds such as

$$\max_{f \in F} \min_{h \in H} \|f - h\|_{\infty} \leq C (\log(|F|))^{-\beta} \quad (13)$$

we have 12 and we can apply 13.

3.2.5 Curse of dimensionality

We look for a bound of the type of 13, where F is the class of uniformly Lipschitz continuous ($\alpha = 1$), and when h is a nearest neighbor algorithm.

We show that in order to have $\sup_{f \in F} \|f - \hat{f}\|_\infty \leq \epsilon$, we must have

$$n \geq \frac{\epsilon^{-d} d^{\frac{d}{2}}}{(2\pi e)^{\frac{d}{2}}}$$

Remark. With α larger, the result is similar.

3.3 Ordinary Least Squares (OLS)

3.3.1 Context

The OLS is an important supervised learning problem. In the least squares problem, the loss l writes

$$l(y, y') = (y - y')^2$$

In the Ordinary Least Squares (OLS) setup, $\mathcal{X} = \mathbb{R}^d$, $\mathcal{Y} = \mathbb{R}$, and the estimator is a **linear function** parametrized by $\theta \in \mathbb{R}^d$.

$$F = \{x \mapsto x^T \theta, \theta \in \mathbb{R}^d\}$$

The dataset is stored in the **design matrix** $X \in \mathbb{R}^{n \times d}$.

$$X = \begin{pmatrix} x_1^T \\ \vdots \\ x_i^T \\ \vdots \\ x_n^T \end{pmatrix} = \begin{pmatrix} x_{11}, \dots, x_{1j}, \dots, x_{1d} \\ \vdots \\ x_{i1}, \dots, x_{ij}, \dots, x_{id} \\ \vdots \\ x_{n1}, \dots, x_{nj}, \dots, x_{nd} \end{pmatrix}$$

The vector of predictions of the estimator writes $Y = X\theta$. Hence,

$$\begin{aligned} R_n(\theta) &= \frac{1}{n} \sum_{i=1}^n (y_i - \theta^T x_i)^2 \\ &= \frac{1}{n} \|Y - X\theta\|_2^2 \end{aligned}$$

Remark. A **linear model**, such as the OLS, can often be interpreted as predicting an output value (dependent variable) from combining the contributions from the **d features** (independent variables) of the input data, in a linear way. This can be useful for classification as well as regression.

3.3.2 OLS estimator

We assume that X is **injective**. Necessary, $d \leq n$.

Proposition 14. Closed form solution

We X is injective, there exists a unique minimiser of $R_n(\theta)$, called the **OLS estimator**, given by

$$\hat{\theta} = (X^T X)^{-1} X^T Y \quad (14)$$

Remark. We consider the non-centered empirical covariance matrix :

$$\hat{\Sigma} = \frac{1}{n} X^T X \in \mathbb{R}^{d,d} \quad (15)$$

Then,

$$\hat{\theta} = \frac{1}{n} \hat{\Sigma}^{-1} X^T Y \quad (16)$$

Démonstration. $R_n(\theta)$ is coercive, continuous, and in a finite dimensional space. Hence, it admits at least a minimizer. It is also differentiable and convex, so the minimizers are the points cancelling the gradient.

Let $u : x \mapsto \|u\|_2^2$. If du_x is the differential of u in x , for all $x \in \mathbb{R}^d$, we have

$$du_x(h) = 2\langle x, h \rangle$$

Let $dR_n(\theta)$ be the differential of R_n in θ .

$$\begin{aligned} dR_n(\theta)(h) &= \frac{1}{n} du_{Y-X\theta}(-Xh) \\ &= \frac{2}{n} \langle Y - X\theta, -Xh \rangle \\ &= \frac{2}{n} \langle X^T(X\theta - Y), h \rangle \end{aligned}$$

where we have used 1.1.1. We deduce that the gradient of R_n in θ writes

$$\nabla_{\theta} R_n = \frac{2}{n} (X^T X \theta - X^T Y) \quad (17)$$

Let us show that $X^T X$ is **invertible**. As it is a square matrix $\in \mathbb{R}^{d \times d}$, we only need to show that it is injective. Let $u \in \mathbb{R}^d$, such that $X^T X u = 0$. Then $\langle X^T X u, u \rangle = 0$, but this inner product equals $\langle Xu, Xu \rangle = \|Xu\|^2$. Hence $Xu = 0$, but as X is injective, $u = 0$, which completes the proof.

We conclude that the only solution to 17 is

$$\hat{\theta} = (X^T X)^{-1} X^T Y$$

□

Remark. We deduce that $\hat{Y} = X\hat{\theta} = X(X^T X)^{-1} X^T Y$ is the orthogonal projection of Y on the linear subspace $\text{Im}(X)$. Hence $P_X = X(X^T X)^{-1} X^T$ is the projection matrix on $\text{Im}(X)$.

Remark. Numerical resolution

The cost of inverting $X^T X$ is $\mathcal{O}(d^3)$ by the Gauss-Jordan method and is thus prohibitive in high dimensions (for instance $> 10^5$).

3.3.3 Statistical analysis of OLS

We are interested in quantitative guarantees on the generalization error of the OLS estimator. In order to compute such guarantees, it is necessary to make some assumptions about the way the data are generated.

Definition 28. Linear model

The assumption that there exists a vector $\theta^* \in \mathbb{R}^d$ such that

$$Y_i = \theta^{*T} x_i + Z_i, \forall i \in [1, n]$$

and Z_i is a centered noise (or error) ($E[Z_i] = 0$) with variance σ^2 , is called the **linear model**.

All expectations are then with respect to the law of X and of Z .

Remark. When $\sigma^2 = 0$, we are in the *noiseless model*.

Definition 29. Fixed design

In this setting, X is assumed to be deterministic, so the expectations are with respect to Z only. We can thus define

$$R_X(\theta) = E\left[\frac{1}{n} \sum_{i=1}^n (Y_i - \theta^T x_i)^2\right]$$

For instance $R_X(\theta^*) = \sigma^2$.

Definition 30. Random design

In this setting, X is random, so the expectations are with respect to X and Z .

Comment : We will give results for the **fixed design** setting in the linear model. The fixed design setting is a **simplifying assumption**. In that respect, it allows to have guarantees on the **fixed design risk** $R_X(\theta)$, not the generalization error directly, with

$$R_X(\theta) = \mathbb{E}_y \left[\frac{1}{n} \|y - X\theta\|_2^2 \right] \quad (18)$$

where the expectation is only over y , and the index X noting that this depends on the input matrix X . The fixed design setting can also be seen as a method that allows the learning of the optimal prediction vector, given some inputs.

However, it is also possible to have similar results in the random design setting.

Proposition 15. Risk decomposition

Let $\|\theta\|_{\hat{\Sigma}}^2 = \theta^\top \hat{\Sigma} \theta$ (it is called the Mahalanobis distance norm). Under the linear model with fixed design, for any fixed $\theta \in \mathbb{R}^d$, we have :

$$R_X(\theta) - R_X(\theta^*) = \|\theta - \theta^*\|_{\hat{\Sigma}}^2 \quad (19)$$

and, if θ is a random variable,

$$\mathbb{E}[R_X(\theta)] - R_X(\theta^*) = \|\mathbb{E}[\theta] - \theta^*\|_{\hat{\Sigma}}^2 + \mathbb{E} \left[\|\theta - \mathbb{E}[\theta]\|_{\hat{\Sigma}}^2 \right] \quad (20)$$

— $\|\mathbb{E}[\theta] - \theta^*\|_{\hat{\Sigma}}^2$: Bias term

— $\mathbb{E} \left[\|\theta - \mathbb{E}[\theta]\|_{\hat{\Sigma}}^2 \right]$: Variance term

Proposition 16. Statistical property of OLS estimator

Under the same hypothesis (linear model, fixed design), the OLS estimator $\hat{\theta}$ defined in 14 satisfies :

— $\hat{\theta}$ is **unbiased** : $\mathbb{E}[\hat{\theta}] = \theta^*$.

— $\text{Var}(\hat{\theta}) = \frac{\sigma^2}{n} \hat{\Sigma}^{-1}$.

$\hat{\Sigma}^{-1}$ is often called the *precision matrix*.

Corollary 4. Distance to optimal parameter, excess risk of OLS

Still with the same hypothesis

$$\text{Var}(\|\hat{\theta} - \theta^*\|^2) = \frac{d\sigma^2}{n} \text{Tr}(\hat{\Sigma}^{-1})$$

and

$$\mathbb{E}[R_X(\hat{\theta})] - R_X(\theta^*) = \frac{\sigma^2 d}{n}$$

We note that both these quantities increase linearly with the dimension.

This corollary is illustrated in figure 1.

Proposition 17. The expected value of the empirical risk of $\hat{\theta}$ writes :

$$\mathbb{E}[R_n(\hat{\theta})] = \frac{n-d}{n} \sigma^2 \quad (21)$$

Démonstration. Not stated here, the proof is obtained by direct calculations and useful matricial tricks. \square

Remark. In this expected value, both the dataset and $\hat{\theta}$ are random variables.




FIGURE 1 – Dependence of the excess risk on n and d , in the setting used in TP1.

3.4 Ridge regression

3.4.1 Context

When the dimension d becomes larger (i.e. $\frac{d}{n}$ approaches 1 or becomes larger), several undesirable phenomena can appear :

- the excess risk is of order σ^2 .
- the training data can be perfectly fit, and this often leads to a bad generalization performance.
- if $d > n$, $X^T X$ is not invertible and there is a linear subspace of points canceling the gradient of the objective function. The problem is said to be **poorly posed** or **unidentifiable**.

Regularizing the problem is an approach to enforce the unicity of the solution at the cost of introducing a **bias** in the estimator. The unicity is guaranteed by the **strong convexity** of the new loss function.

3.4.2 Ridge regression estimator

Definition 31. Ridge regression estimator

It is defined as

$$\hat{\theta}_\lambda = \arg \min_{\theta \in \mathbb{R}^d} \left(\frac{1}{n} \|Y - X\theta\|_2^2 + \lambda \|\theta\|_2^2 \right) \quad (22)$$

Proposition 18. The Ridge regression estimator is unique even if $X^T X$ is not invertible and is given by

$$\hat{\theta}_\lambda = \frac{1}{n} (\hat{\Sigma} + \lambda I_d)^{-1} X^T Y$$

Démonstration. We can show that the loss 22 is strongly convex. To do so, it is sufficient to show that $x \mapsto \|\theta\|^2$ is 2-convex on \mathbb{R}^d . We have

$$\forall u, v \in \mathbb{R}, \left(\frac{u+v}{2} \right)^2 = \frac{u^2}{2} + \frac{v^2}{2} - \frac{1}{4}(u-v)^2$$

Hence in \mathbb{R} , $u \mapsto u^2$ is 2-convex. By summing linear mappings (the projections $\theta \mapsto \theta_i$) and 2-convex applications ($\theta_i \mapsto \theta_i^2$), $\theta \mapsto \|\theta\|^2$ is 2-convex in \mathbb{R}^n .

By summation with $\theta \mapsto \frac{1}{n}\|Y - X\theta\|$, which is a convex function, $\theta \mapsto \tilde{R}_n(\theta)$ is 2λ -convex.

Hence there exists a unique minimizer for R_n . As it is differentiable, the minimizer is uniquely defined by cancellation of the gradient. The gradient in θ writes

$$\frac{2}{n}(X^T X \theta - X^T y) + 2\lambda \theta$$

The equation of the cancellation of the gradient is

$$\left(\frac{2}{n}n\hat{\Sigma} + 2\lambda I_d\right)\theta_\lambda = \frac{2}{n}X^T y$$

which we can write

$$n(\hat{\Sigma} + \lambda I_d)\theta_\lambda = X^T y$$

$\hat{\Sigma} + 2\lambda I_d$ is a symmetric matrix with all eigenvalues $\geq 2\lambda$. Thus, it is invertible. Also, $\forall a \in \mathbb{R}^*$ and $A \in GL_d(\mathbb{R})$, $(aA)^{-1} = \frac{1}{a}A^{-1}$, which concludes the proof. \square

3.4.3 Statistical analysis of Ridge regression

We can also analyse the statistical properties of the ridge estimator, with the same hypotheses as before : linear model, fixed design. It is also possible to do it in the random design and potentially infinite d .

Proposition 19. *Under the linear model assumption, with fixed design setting, the ridge regression estimator has the following excess risk*

$$\mathbb{E}[R(\hat{\theta}_\lambda)] - R^* = \lambda^2 \theta^{*T} (\hat{\Sigma} + \lambda I_d)^{-2} \hat{\Sigma} \theta^* + \frac{\sigma^2}{n} \text{tr}[\hat{\Sigma}^2 (\hat{\Sigma} + \lambda I_d)^{-2}] \quad (23)$$

Démonstration. As for OLS, the proof is obtained by direct calculation and useful matricial tricks. \square

Comments :

- We observe again a bias / variance decomposition.
- We consider the bias term B :

$$B = \lambda^2 \theta^{*T} (\hat{\Sigma} + \lambda I_d)^{-2} \hat{\Sigma} \theta^* \quad (24)$$

- The bias B increases when λ increases. It is an approximation error and does not depend on n .
- When $\lambda = 0$ and $\hat{\Sigma}$ is invertible (which corresponds to OLS), $B = 0$.
- When $\lambda \rightarrow +\infty$, $B \rightarrow \theta^{*T} \hat{\Sigma} \theta^*$.
- We consider the variance term V :

$$V = \frac{\sigma^2}{n} \text{tr}[\hat{\Sigma}^2 (\hat{\Sigma} + \lambda I_d)^{-2}] \quad (25)$$

- The variance V decreases when λ increases. It is an estimation error and depends on n .
- When $\lambda = 0$ and $\hat{\Sigma}$ is invertible (which corresponds to OLS), $V = \frac{\sigma^2 d}{n}$.
- When $\lambda \rightarrow +\infty$, $V \rightarrow 0$.
- When $n \rightarrow +\infty$, $V \rightarrow 0$.
- In most cases, it is preferable to have a biased estimation ($\lambda > 0$).

3.4.4 Choice of λ

A natural question is whether it is possible to have a lower excess risk with Ridge regression than with OLS, which means an excess risk smaller than $\frac{\sigma^2 d}{n}$.

Proposition 20. *With the choice*

$$\lambda^* = \frac{\sigma \sqrt{\text{tr}(\hat{\Sigma})}}{\|\theta^*\|_2 \sqrt{n}} \quad (26)$$

then

$$\mathbb{E}[\mathcal{R}(\hat{\theta}_\lambda)] - \mathcal{R}^* \leq \frac{\sigma \sqrt{\text{tr}(\hat{\Sigma})} \|\theta^*\|_2}{\sqrt{n}} \quad (27)$$

This proposition is illustrated in figure 2.

- If the norms of all x_i are uniformly bounded by R , then we can observe by direct calculation that $\text{tr}(\Sigma) \leq R^2$, independently of d . If R and $\|\theta^*\|_2$ remain finite, then d plays no role in the bound in 27. This is called a *dimension-free* bound.
- The convergence to 0 in OLS is in $\frac{1}{n}$, while it is in $\frac{1}{\sqrt{n}}$ for the ridge. However, for the ridge regression, the dependence in the noise is in σ , whereas it is σ^2 for OLS. Which one is preferable will depend on the value of the constants, and will not necessary be the "fast" rate in $\mathcal{O}(\frac{1}{n})$.
- In practical situations, the quantities involved in the computation of λ^* in 26 are typically unknown. However this equation show that there may exist a λ with a good prediction performance, which can be found by cross validation in practice.




FIGURE 2 – Comparison of OLS and Ridge performance for a specific setting used in TP2.

3.5 Lasso

3.5.1 Context

The **Lasso** is another regularization method for OLS. The original idea is to penalize the number of nonzero components of the found parameter θ . However, if we were to do it using a penalization of the form $\lambda\|\theta\|_0$ ($\|\theta\|_0$ is precisely the number of nonzero components), this would lead to a non-convex optimization problem. Lasso replaces $\|\theta\|_0$ by $\|\theta\|_1$.

3.5.2 LASSO estimator

It is defined as

$$\tilde{\theta}_\lambda \in \arg \min_{\theta \in \mathbb{R}^d} \{\|Y - X\theta\|^2 + \lambda\|\theta\|_1\}$$

3.5.3 Numerical resolutions

The problem is convex, so the solution can be obtained efficiently. Popular methods include **Coordinate descent**, **Fista**, **LARS**.

3.6 Logistic regression

Despite its name, **logistic regression** is a classification method. It is sometimes called logistic classification.

3.6.1 Context

We want to perform a binary classification from d -dimensional inputs, using a **linear predictor** \hat{f}_θ , parametrized by $\theta \in \mathbb{R}^d$. The inputs are $x_i \in \mathbb{R}^d$, $1 \leq i \leq n$ and the outputs : $y_i \in \{0, 1\}$. The estimator writes :

$$\hat{y}_i = \hat{f}_\theta(x_i) = \begin{cases} 1 & \text{if } \theta^T x_i \geq 0 \\ 0 & \text{if } \theta^T x_i < 0 \end{cases}$$

Remark. *It is not always possible to separate the data in a linear way. Often we work with **transformations** of the data, $\phi(x)$ instead of x , where $\phi(x)$ is in a high dimensional space (Kernel methods).*

Definition 32. Binary loss function

$$\hat{R}_n(\theta) = \frac{1}{n} \sum_{i=1}^n \mathbb{1}_{y_i \neq \hat{f}_\theta(x_i)}$$

However it is hard to minimize the binary loss as it is neither differentiable nor convex in θ . We can replace it by a **convex, differentiable surrogate loss (substitut convexe)**. Several possibilities exist instead of using $\mathbb{1}_{y_i \neq \hat{f}_\theta(x_i)}$ as l (binary loss). The **logistic loss** is one of them

Definition 33. Logistic loss

$$l(\hat{y}, y) = y \log(1 + e^{-\hat{y}}) + (1 - y) \log(1 + e^{\hat{y}}) \quad (28)$$

Remark. *If $Y \in \{-1, 1\}$, we sometimes encounter the following formulation :*

$$l(\hat{y}, y) = \log(1 + e^{-\hat{y} \times y})$$

Another often used possibility is the **Hinge loss**.

3.6.2 Logistic regression estimator

Definition 34. Logistic regression estimator

If l is the logistic loss, it is defined as

$$\hat{\theta}_{\text{logit}} = \arg \min_{\theta \in \mathbb{R}^d} \frac{1}{n} \sum_{i=1}^n l(X_i^T \theta, Y_i)$$

Numerical resolution : The ER is differentiable and strictly convex and thus has a unique minimizer. However, the equation obtained by cancellation of the gradient has **no closed form solution**. To solve it, we need to use iterative algorithms such as gradient descent or Newton's method.

Regularization : To avoid overfitting, we may also add regularization to the logistic loss, for instance with a term $\lambda \|\theta\|_2^2$.

3.7 Stochastic Gradient Descent (SGD)

3.7.1 Context

The SGD is the standard optimization algorithm for large scale machine learning, as its computational complexity is lower than that of GD, while keeping good convergence speed. We want to minimise a function of the form

$$f(\theta) = \mathbb{E}_z L(\theta, z)$$

where

$$\mathbb{E}_z L(\theta, z) = \int_{\Omega} L(\theta, z) d\rho(z)$$

- $z \in \Omega$, ρ is a probability distribution over Ω .
- $L : \mathbb{R}^d \times \Omega \Rightarrow \mathbb{R}$
- $\theta \in \mathbb{R}^d$ is the parameter we want to optimize.
- for instance $z = (x_i, y_i)$ as in supervised learning and

$$L(\theta, z) = l(\theta^T x, y)$$

The idea of SGD is to use at each time t an **unbiased estimator** g_t of the gradient of f , instead of computing the full gradient. Indeed, for large problems, this can be prohibitive computationally. Hence, for large-scale machine learning, SGD and its variants are the most widely used optimization algorithms. Let us discuss this on two examples.

1] Empirical risk minimisation (ERM)

For instance, f can be the empirical risk f_n .

$$f_n(\theta) = \frac{1}{n} \sum_{i=1}^n L(\theta, z_i)$$

By **linearity** of the expected value,

$$\nabla_{\theta} f_n = \frac{1}{n} \sum_{i=1}^n \nabla_{\theta} L(\theta, z_i) \quad (29)$$

If at time t we sample uniformly $i(t) \in \{1, \dots, n\}$, then if we define g_t as

$$g_t = \nabla_{\theta} L(\theta, z_{i(t)})$$

then g_t is an unbiased estimator of $\nabla_{\theta} f_n$. SGD uses g_t instead of the **batch gradient** $\nabla_{\theta} f_n$.

2] Expected risk minimization

If the hypotheses of Leibniz theorem are satisfied (which is the case most of the time in Machine Learning), we can write

$$\begin{aligned}\nabla_{\theta} E[L(\theta, z)] &= \nabla_{\theta} \int_{\Omega} L(\theta, z) dz \\ &= \int_{\Omega} \nabla_{\theta} L(\theta, z) dz \\ &= E[\nabla_{\theta} L(\theta, z)]\end{aligned}$$

Hence, if $z \sim \rho$, $\nabla_{\theta} L(\theta, z)$ is an **unbiased estimator** of the gradient of the **real risk**.

However, importantly, in real situations, we can not sample infinitely many $z \sim \rho$. We only have access to z_{i_t} , sampled from the finite dataset D_n , and thus $z_{i_t} \sim \hat{\rho}_n$.

3.7.2 SGD updates

The SGD update reads

$$\theta_t = \theta_{t-1} - \gamma_t \nabla_{\theta} L(\theta, z_{i_t})$$

where z_{i_t} is the sample obtained at time t .

3.7.3 SGD as an estimation of GD

Let us compute the expected value of the SGD update over z_{i_t} sampled from ρ , given a learning rate γ .

$$\begin{aligned}E_{z_{i_t}}(\theta_t) &= \theta_{t-1} - \gamma E_{z_{i_t}} \nabla_{\theta_{t-1}} L(\theta_{t-1}, z_{i_t}) \\ &= \theta_{t-1} - \gamma \nabla_{\theta} f(\theta_{t-1})\end{aligned}$$

Hence, in expectation, SGD behaves as GD.

3.7.4 Convergence for a convex objective

Reference : [Bach, 2021,]

Theorem 13. *Convergence of SGD for a convex function*

We assume that :

- The gradients estimates are unbiased.
- f is convex
- f is L -Lipshitz
- The gradient is bounded : $\forall t, \|g_t(\theta_{t-1})\|_2^2 \leq L^2$ almost surely.
- f admits a minimiser θ^* such that $\|\theta^* - \theta_0\|_2 \leq D$.

Let $\gamma_t = \frac{D}{L\sqrt{t}}$, $\forall t \in \mathbb{N}$. Then, the iterates of SGD satisfy :

$$E[f(\hat{\theta}_t) - f(\theta^*)] \leq DL \frac{2 + \log(t)}{2\sqrt{t}}$$

where

$$\hat{\theta}_t = \frac{\sum_{s=1}^t \gamma_s \theta_{s-1}}{\sum_{s=1}^t \gamma_s}$$

3.7.5 Convergence for a strongly convex objective

We replace f by g to obtain a strongly convex objective and consider the minimization of g .

$$g(\theta) = f(\theta) + \frac{\mu}{2} \|\theta\|_2^2$$

Now, the SGD iteration reads :

$$\theta_t = \theta_{t-1} - \gamma_t (g_t(\theta_{t-1}) + \mu \theta_{t-1})$$

Theorem 14. *Convergence of a strongly convex objective*

We use the same hypotheses as 13, adding the assumption that g admits a (necessary unique) minimizer θ^* . Then, with $\gamma_t = \frac{1}{\mu t}$.

$$\mathbb{E}[f(\hat{\theta}_t) - f(\theta^*)] \leq \frac{2L^2(1 + \log t)}{\mu t}$$

where

$$\hat{\theta}_t = \frac{1}{t} \sum_{s=1}^t \theta_{s-1}$$

3.7.6 Convergence for a strongly convex L -smooth objective

Definition 35. Variance of the estimator of the gradient

We define the **variance** of the estimator of the gradient as

$$\sigma^2(\theta) = \mathbb{E}_z \|\nabla_{\theta} f(\theta) - \nabla_{\theta} L(\theta, z)\|^2$$

We note that σ depends on θ .

Theorem 15. *Convergence of SGD for a strongly convex L -smooth function*

We assume that :

- The gradients estimates are unbiased.
- f is μ strongly convex.
- f has L -Lipshitz continuous gradients
- $\exists \sigma^2 > 0, \forall \theta \in \mathbb{R}^d, \sigma^2(\theta) \leq \sigma^2$

Let $\gamma_t = \gamma, \forall t \in \mathbb{N}$ and $0 < \gamma < \frac{1}{2L}$. Then

$$\mathbb{E}_{z_1, \dots, z_T} \|\theta_T - \theta^*\| \leq (1 - \mu\gamma)^T \|\theta_0 - \theta^*\| + \frac{\gamma}{\mu} \sigma^2$$

Remark. The $\log t$ factors at the numerator are often ignored when considering the computational complexities.

3.7.7 Computational complexity of GD and SGD

Roughly speaking, in most situations the computational complexity of one iteration of GD will be n times that of one iteration of SGD. For instance, in the case of least-squares regression, computing a batch gradient (GD) costs $\mathcal{O}(nd)$. Computing a SGD gradient estimate costs $\mathcal{O}(d)$. For example, in the case of Ridge regression, the loss is μ -strongly convex and L -smooth, with a condition number $\kappa = \frac{L}{\mu}$. If t denotes the number of iterations, then

- GD has a convergence rate of $\mathcal{O}(\exp(-\frac{t}{\kappa}))$. To get an error of ϵ , we must have $t = \mathcal{O}(\kappa \log \frac{1}{\epsilon})$. Since each iteration requires $\mathcal{O}(nd)$ computations, the computation time will be $\mathcal{O}(\kappa nd \log \frac{1}{\epsilon})$.
- SGD has a convergence rate of $\mathcal{O}(\frac{\kappa}{t})$. To get an error of ϵ , we must have $t = \mathcal{O}(\frac{\kappa}{\epsilon})$. Since each iteration is $\mathcal{O}(d)$, we have a computation time of $\mathcal{O}(\frac{\kappa d}{\epsilon})$.

As a consequence :

- When n is large and ϵ not too small, GD will need more computation time to reach error ϵ . An order of magnitude can be obtained by studying the value ϵ^* such that

$$\kappa n d \log \frac{1}{\epsilon^*} = \frac{\kappa d}{\epsilon^*}$$

Which translates to

$$\epsilon^* \log \epsilon^* = -\frac{1}{n}$$

- When $\epsilon \rightarrow 0$, GD becomes faster than SGD to reach this precision.

To conclude, for low precision and large n , SGD is a preferable. Also, in machine learning, due to the estimation error that is $\mathcal{O}(\frac{1}{\sqrt{n}})$, a very high precision is often not needed. Improvements of the method exist, such as variance reduction methods, SAG or SAGA [Schmidt et al., 2013,].

3.8 Kernels

3.8.1 Context

Kernel methods replace inputs x by a function $\phi(x) \in \mathbb{R}^d$ and then perform linear predictions on $\phi(x) : f(x) = \theta^T \phi(x)$. $\phi(x)$ is called the **feature** associated to x . The reason why they are called "kernel" is not obvious at first glance but will be justified after some calculations.

Examples :

- **Polynomial regression of degree d .** $x \in \mathbb{R}$, $\phi(x) = (1, x, x^2, \dots, x^d)$. If $\theta \in \mathbb{R}^d$,

$$f(x) = \theta^T \phi(x) = \sum_{j=1}^r \theta_j x^j$$

- **Polynomial multivariate regression of degree r .** $x \in \mathbb{R}^p$.

$$\phi(x) = \{(x_1^{\alpha_1} \dots x_p^{\alpha_p}), \sum_{i=1}^p \alpha_i = r\}$$

In that case, $d = \binom{d+r-1}{r}$.

- **Fourier Basis.** $\phi_1, \dots, \phi_d : \mathbb{R}^p \rightarrow \mathbb{R}$ with $\phi_k(x) = e^{i \frac{2\pi k x}{T}}$.

$$f(x) = \sum_{i=1}^d w_i \phi_i(x)$$

3.8.2 Representer theorem

We place ourselves in a general framework where we want to study the minimizer $\hat{\theta}$ of a loss such as

$$L(w) = \frac{1}{n} \sum_{i=1}^n \ell(y_i, \theta^T \phi(x_i)) + \lambda \|\theta\|^2 \quad (30)$$

With $\lambda > 0$ and $\phi(x_i) \in \mathbb{R}^d$. The following theorem states that the minimizer belongs to the span of $\phi(x_i)$ (espace engendré), which is a subspace of \mathbb{R}^d . This will have important consequences.

Theorem 16. *Representer theorem*

Let $\Psi : \mathbb{R}^{n+1} \rightarrow \mathbb{R}$ be a strictly increasing function with respect to the last variable. Then, the minimum of

$$L(\theta) = \Psi(\theta^T \phi(x_1), \dots, \theta^T \phi(x_n), \|\theta\|^2)$$

is attained for $\hat{\theta} \in \text{Vect}(\{\phi(x_i)\})$. We can write

$$\theta = \sum_{i=1}^n \alpha_i \phi(x_i), \alpha_i \in \mathbb{R}$$

Corollary 5. As a direct consequence, the minimum of loss 30 is attained at $\theta = \sum_{i=1}^n \alpha_i \phi(x_i)$, $\alpha_i \in \mathbb{R}$.

We note that no convexity hypothesis on l is required.

3.8.3 Consequence

We note

- α the vector such that $\theta = \sum_{i=1}^n \alpha_i \phi(x_i)$.
- $K \in \mathbb{R}^{n,n}$ the matrix defined by

$$K_{ij} = \phi(x_i)^T \phi(x_j) = \langle \phi(x_i), \phi(x_j) \rangle$$

Now, we remark that

$$\begin{aligned} \|\theta\|^2 &= \langle \theta, \theta \rangle \\ &= \left\langle \sum_{i=1}^n \alpha_i \phi(x_i), \sum_{j=1}^n \alpha_j \phi(x_j) \right\rangle \\ &= \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j \langle \phi(x_i), \phi(x_j) \rangle \\ &= \sum_{i=1}^n \alpha_i \left(\sum_{j=1}^n K_{ij} \alpha_j \right) \\ &= \sum_{i=1}^n \alpha_i (K\alpha)_i \\ &= \alpha^T K \alpha \end{aligned}$$

And that $\forall i \in [1, n]$,

$$\begin{aligned} \theta^T \phi(x_i) &= \sum_{j=1}^n \alpha_j \phi(x_j)^T \phi(x_i) \\ &= \sum_{j=1}^n \alpha_j K_{ij} \\ &= (K\alpha)_i \end{aligned}$$

Finally

$$L(\theta) = \frac{1}{n} \sum_{i=1}^n l(y_i, (K\alpha)_i) + \lambda \alpha^T K \alpha$$

Hence, the loss can be written **only with** K and α , instead of the explicit $\phi(x_i)$. At first glance, this could seem like a deceitful statement, since we should need to know $\phi(x_i)$ and $\phi(x_j)$ in order to compute $k(x_i, x_j)$. **However**, in some situations, it is possible to compute $k(x_i, x_j)$ **without explicit knowledge** of ϕ . This is known as the **kernel trick**.

More generally, we define **kernel functions**.

Definition 36. Kernel function

Let \mathcal{X} be a non-empty set. A function $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ is called a **positive-definite kernel** if there exists a \mathbb{R} -Hilbert space \mathcal{H} , with an inner product $\langle \cdot, \cdot \rangle_{\mathcal{H}}$, and a map $\phi : \mathcal{X} \rightarrow \mathcal{H}$ such that $\forall x, x' \in \mathcal{X}$,

$$k(x, x') = \langle \phi(x), \phi(x') \rangle_{\mathcal{H}}$$

Remark. No hypothesis is made on \mathcal{X} .

We come back to our example. We can rewrite the evaluation function as :

$$f(x) = \theta^T \phi(x) = \sum_{i=1}^n \alpha_i k(x_i, x)$$

Consequences :

- The initial optimization problem on θ can be replaced as an optimization problem on α .
- If $d \gg n$ (or even if d is infinite, which can happen as we will see below), and if it is relatively fast to compute k , this can be a computational improvement. However, computing the matrix K requires at least $\mathcal{O}(n^2)$ operations.

Example : here is an example where we can compute $k(x, x')$ without computing $\phi(x)$ explicitly. We consider $x \in]-1, 1[$, and

$$\phi(x) = (1, x, x^2, \dots)$$

Then,

$$\begin{aligned} k(x, x') &= \langle (1, x, x^2, \dots), (1, x', x'^2, \dots) \rangle \\ &= \sum_{i \in \mathbb{N}^*} x^i x'^i \\ &= \sum_{i \in \mathbb{N}^*} (xx')^i \\ &= \frac{1}{1 - xx'} \end{aligned}$$

Although $\phi(x)$ is **infinite dimensional**, we can compute $k(x, x')$ without evaluating and storing $\phi(x)$ in memory (which would be impossible anyways).

4 UNSUPERVISED LEARNING

4.1 Definitions

4.1.1 Setup

In an **Unsupervised Learning** context, we are given a dataset $\{x_i, i \in [1, n]\}$, from which we want to extract information. Most of the time the x_i are vectors : $\forall i, x_i \in \mathbb{R}^d$. Most common applications include :

- **Clustering** : identifying **groups** in the dataset by forming **partition**.
- **Dimensionality reduction** : transforming the data x_i into points z_i , most of time in a vector space of smaller dimension than the original space. The z_i can then be used with less computational burden by a supervised learning algorithm, for instance.
- **Density estimation** : computing a useful approximation of the probability law that generated the dataset.

4.2 K-means clustering

4.2.1 Context

K-means clustering is a **vector quantization** method. Vector quantization associates **prototypes**, or **centroids**, to the data. Given a centroid $\mu_k \in \mathbb{R}^p$, the cluster k is the set of all points for which the closest centroid is μ_k .

The centroids μ_k are stored in a matrix μ , the z_i^k in matrix z .

z_i^k is the indicator variable associated to x_i . If x_i belongs to cluster k , $z_i^k = 1$. Otherwise $z_i^k = 0$.

4.2.2 Distorsion

Definition 37. Distorsion

$$J(\mu, z) = \sum_{i=1}^n \sum_{k=1}^K z_i^k \|x_i - \mu_k\|^2$$

The K-means minimizes the distorsion by **Expectatoin-maximization algorithm**.

4.3 Principal Component Analysis (PCA)

4.3.1 Context

The PCA is a linear dimension reduction technique. Points in \mathbb{R}^d are linearly projected on a well chosen affine subspace.

4.3.2 Optimal projection

Without loss of generality, we assume the data are **centered**. $\sum_{i=1}^n x_i = 0$. X is the design matrix as in 3.3.1.

We look for $w \in \mathbb{R}^d$, with $\|w\| = 1$, such that $\hat{\text{Var}}(w^T x)$ is maximal. Here, $\hat{\text{Var}}$ denotes the empirical variance (see 1.2.2).

Proposition 21. w is the eigenvector of $X^T X$ with largest eigenvalue λ_{\max} .

Démonstration. We want to maximize $\hat{\text{Var}}(w^T x)$, which is equivalent to maximizing

$$\begin{aligned} \sum_{i=1}^n (w^T x_i)^2 &= w^T (X^T X) w \\ &= \langle (X^T X) w, w \rangle \leq \lambda_{\max} \end{aligned}$$

□

Another point of view on the problem is to define w as

$$w = \arg \min_{\|w\|=1} \sum_{i=1}^n d(x_i, \{\lambda w, \lambda \in \mathbb{B}\})^2$$

4.3.3 Algorithm

Most of the time we project on a **subspace**, not only on a vector. The sequence of vectors obtained are the first normalized eigenvectors of $X^T X$, which are orthogonal.

4.4 Density estimation

4.4.1 Context

From points $(x_1, \dots, x_n) \in \mathbb{R}^d$, we would like to estimate their unknown distribution probability ρ . What we have access to is the empirical distribution ρ_n ,

$$\rho_n(x) = \frac{1}{n} \sum_{i=1}^n \delta_{x_i}$$

Where δ_{x_i} is the Dirac distribution in x_i .

4.4.2 Kernel density estimation (KDE) and Convolution

Let q be a probability density on \mathbb{R}^d , for instance, as $q(x) = e^{-\|x\|^2}$. Let $\tau > 0$. We define a scaling of q as

$$q_\tau(x) = \tau^{-d} q\left(\frac{x}{\tau}\right)$$

We verify that q_τ is also a probability density.

$$\begin{aligned} \int_{\mathbb{R}^d} q_\tau(x) dx &= \int_{\mathbb{R}^d} \tau^{-d} q\left(\frac{x}{\tau}\right) dx \\ &= 1 \end{aligned}$$

We define the estimator $\hat{\rho}_\tau$ as

$$\hat{\rho}(x) = \frac{1}{n} \sum_{i=1}^n q_\tau(x - x_i)$$

This estimation method is called **Kernel density estimation (KDE)**, or **Parzen method**. If $*$ is the convolution of f and g , $(f * g)(x) = \int f(y)g(x - y)dy$, then

$$\hat{\rho}(x) = \rho_n * q_\tau$$

5 PROBABILISTIC MODELING

5.1 Context

We are given a set of observations $\{y_1, \dots, y_n\} \in \mathcal{Y}$ that we assume are generated i.i.d from an unknown distribution. We look for a **probabilistic model** that explains well the data. We could use this model to predict well new data, that would be statistically similar to the observed ones.

Definition 38. Parametric model

Let μ be a measure on \mathcal{Y} . Most of time times in Machine Learning, μ is the counting measure if $\mathcal{R} \subset \mathbb{N}$ and the Lebesgue measure if $\mathcal{Y} \subset \mathbb{R}^d$.

Let $d > 1$ and $\Theta \subset \mathbb{R}^p$ be a set of parameters. A parametric model \mathcal{P} is a set of probability distributions on \mathcal{Y} with a density with respect to μ , indexed by Θ .

$$\mathcal{P} = \{p_\theta d\mu | \theta \in \Theta\}$$

If we assume that the data were generated from some $p_{\theta^*} \in \mathcal{P}$, with a unknown parameter θ^* , our goal is to find a good estimation of θ . If the data are indeed generated by a distribution in \mathcal{P} , the problem is said to be **well specified**. Otherwise, the problem is said to be **misspecified**.

5.2 Maximum likelihood estimation

Definition 39. Likelihood

Let $\mathcal{P} = \{p_\theta, \theta \in \Theta\}$ be a parametric model. Given $y \in \mathcal{Y}$, the **likelihood** of θ is defined as the function $\theta \mapsto p_\theta(y)$.

The likelihood $L(\cdot|D_n)$ of a dataset $D_n = (y_1, \dots, y_n)$ is defined as

$$L(\cdot|D_n) : \theta \mapsto \prod_{i=1}^n p_\theta(y_i)$$

The **maximum likelihood estimator** (MLE) is the parameter θ that maximises the likelihood :

$$\hat{\theta}_n \in \arg \max_{\theta \in \Theta} (L(\theta|D_n))$$

Remark. — Since the samples y_i are assumed to be independent, the likelihood corresponds to the probability of observing the dataset according to p_θ .

— We often maximise the log of the likelihood, as it is easier to differentiate a sum. Since log is an increasing function, the MLE is also the maximiser of the log of L.

5.2.1 Link with empirical risk minimization

In the context of density estimation, we can define a loss function as the **negative log-likelihood**.

$$\Theta \times \mathcal{Y} \mapsto -\log(p_\theta(y))$$

Given this loss, the risk writes :

$$R(\theta) = E_Y[-\log(p_\theta(y))]$$

and the empirical risk (ER) :

$$R_n(\theta) = -\frac{1}{n} \sum_{i=1}^n \log(p_\theta(y_i))$$

The MLE is then also the empirical risk minimizer.

5.2.2 Link with Kullback-Leibler divergence

The Kullback-Leibler divergence is a quantity used to compare two probability distributions.

Definition 40. Kullback-Leibler divergence

Given two distributions p and q , the KL divergence from p to q is defined as :

$$KL(p||q) = E_{Y \sim p}[\log \frac{p(Y)}{q(Y)}]$$

Lemme 10. If the data are generated by p_{θ^*} , then $KL(p_{\theta^*}||p_\theta)$ is the excess risk of p_θ , with the negative log-likelihood loss.

5.2.3 Conditional modeling

The likelihood can also be used to estimate the density of an output Y given an output X . We then have a collection of densities on Y , for each θ and for each x . We also define a risk and an empirical risk the log-likelihood loss (also called **conditional log-likelihood**), as :

$$R(\theta) = E[-\log(p_\theta(Y|X))]$$

$$R_n(\theta) = -\frac{1}{n} \sum_{i=1}^n \log(p_\theta(y_i|x_i))$$

Link with linear regression. We consider $D_n = \{(x_1, y_1), \dots, (x_n, y_n)\} \in (\mathcal{X} \times \mathcal{Y})^n$, with $\mathcal{X} = \mathbb{R}^d$, and $\mathcal{Y} = \mathbb{R}$. We can define a probabilistic model such that finding the maximum likelihood estimator for this model is equivalent to doing least square regression.

This model, called the **Gaussian model**, assumes that the outputs y_i are independently generated from a Gaussian distribution of mean $w_*^T x_i$ and variance σ_* . Our optimal parameter is thus $\theta^* = (w_*, \sigma_*)$. This writes :

$$Y = w_*^T X + \epsilon$$

where $\epsilon \sim \mathcal{N}(0, \sigma_*^2)$. We want an estimation of θ_* . For a parameter $\theta = (w, \sigma)$, the conditional density is :

$$p_\theta(y|x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(y-w^T x)^2}{2\sigma^2}}$$

Hence

$$\begin{aligned} R_n(\theta) &= -\frac{1}{n} \sum_{i=1}^n \log(p_\theta(y_i|x_i)) \\ &= \frac{1}{2n\sigma^2} \sum_{i=1}^n (y_i - w^T x_i)^2 + \frac{1}{2} \log(2\pi\sigma^2) \end{aligned}$$

This means that the maximum likelihood estimator \hat{w}_n in the Gaussian model is the minimiser of the least squares. We can also show that the estimator of $\hat{\sigma}^2$ is

$$\hat{\sigma}_n^2 = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{w}_n^T x_i)^2$$

Link with logistic regression. In a similar way, in a binary classification problem, we can introduce a model for which the maximum likelihood estimator is the logistic regression estimator $\hat{\theta}_{\text{logit}}$ defined in 34. This time, $\mathcal{Y} = \{0, 1\}$, while we still have $\mathcal{X} = \mathbb{R}^d$.

Definition 41. Sigmoid function

$$\sigma : \begin{cases} \mathbb{R} \rightarrow]0, 1[\\ z \mapsto \frac{1}{1+e^{-z}} \end{cases}$$

We have $\forall z \in \mathbb{R}$,

$$\sigma(-z) = 1 - \sigma(z)$$

and

$$\sigma'(z) = \sigma(z)\sigma(-z)$$

Let us now consider the probabilistic model such that

$$p_\theta(1|x) = \sigma(\theta^T x)$$

which makes sense since $\sigma(\theta^T x) \in [0, 1]$, and can thus be interpreted as a probability. Equivalently, this model can be written (remember that $y = 0$ or $y = 1$)

$$p_\theta(y|x) = (\sigma(\theta^T x))^y (1 - \sigma(\theta^T x))^{1-y}$$

Hence

$$\begin{aligned}
R_n(\theta) &= -\frac{1}{n} \sum_{i=1}^n \log(p_\theta(y_i|x_i)) \\
&= -\frac{1}{n} \sum_{i=1}^n \log\left((\sigma(\theta^T x_i))^{y_i} (1 - \sigma(\theta^T x_i))^{1-y_i}\right) \\
&= -\frac{1}{n} \sum_{i=1}^n y_i \log(\sigma(\theta^T x_i)) + (1 - y_i) \log(1 - \sigma(\theta^T x_i)) \\
&= \frac{1}{n} \sum_{i=1}^n y_i \log(1 + e^{\theta^T x_i}) + (1 - y_i) \log(1 - e^{\theta^T x_i}) \\
&= \frac{1}{n} \sum_{i=1}^n l(\theta^T x_i, y_i)
\end{aligned}$$

where l is the logistic loss defined in 33. We find that the problems of empirical risk minimization for the log-likelihood in this model and the logistic regression are the same.

RÉFÉRENCES

- [Allaire, 2012] Allaire, G. (2012). Analyse numérique et optimisation Une introduction à la modélisation mathématique et à la simulation numérique. Éditions de l'École Polytechnique, (2) :480.
- [Bach, 2021] Bach, F. (2021). Learning Theory from First Principles Draft. Book Draft, page 229.
- [Schmidt et al., 2013] Schmidt, M., Le Roux, N., and Bach, F. (2013). Minimizing finite sums with the stochastic average gradient. Mathematical Programming, 162(1-2) :83–112.