


Fondamentaux théoriques du machine learning



Overview of lecture 3

Risks and risk decompositions

- Examples

- Expected value of empirical risk

- Risk decomposition

- Optimization error

Optimization in machine learning

- Existence results

- Convex analysis

- Gradients

Ordinary Least squares II

- OLS estimator

- Statistical analysis of OLS

Risks and risk decompositions

- Examples

- Expected value of empirical risk

- Risk decomposition

- Optimization error

Optimization in machine learning

- Existence results

- Convex analysis

- Gradients

Ordinary Least squares II

- OLS estimator

- Statistical analysis of OLS

Bayes rule

$$P(A \cap B) = P(A|B)P(B) \quad (1)$$

Law of total probability

If for instance $\Omega = A \cup B \cup C$ and A, B, C are mutually exclusive, then

$$P(X) = P(X \cap A) + P(X \cap B) + P(X \cap C) \quad (2)$$

Exercise 1: Consider the following random variable (X, Y) .

► $X \sim B(\frac{1}{2}),$

$$Y = \begin{cases} B(p) & \text{if } X = 1 \\ B(q) & \text{if } X = 0 \end{cases}$$

With $B(p)$ a Bernoulli law with parameter p .

► Hence $\mathcal{X} = \{0, 1\}, \mathcal{Y} = \{0, 1\}.$

Exercise 1: Consider the following random variable (X, Y) .

► $X \sim B(\frac{1}{2}),$

$$Y = \begin{cases} B(p) & \text{if } X = 1 \\ B(q) & \text{if } X = 0 \end{cases}$$

With $B(p)$ a Bernoulli law with parameter p .

► A predictor $f_1 : \{0, 1\} \rightarrow \{0, 1\} :$

$$f_1 = \begin{cases} 1 & \text{if } x = 1 \\ 0 & \text{if } x = 0 \end{cases}$$

With the "0 - 1" loss, what is the risk (generalization error) of f_1 , $R(f_1)$?

Exercise 1: Consider the following random variable (X, Y) .

► $X \sim B(\frac{1}{2}),$

$$Y = \begin{cases} B(p) & \text{if } X = 1 \\ B(q) & \text{if } X = 0 \end{cases}$$

► $f_1 : \{0, 1\} \rightarrow \{0, 1\} :$

$$f = \begin{cases} 1 & \text{if } x = 1 \\ 0 & \text{if } x = 0 \end{cases}$$

$$\begin{aligned} R(f_1) &= E[I(Y, f(X))] \\ &= 1 \times P(Y \neq f(X)) + 0 \times P(Y = f(X)) \\ &= P(Y \neq f(X)) \end{aligned} \tag{3}$$

► $X \sim B(\frac{1}{2}),$

$$Y = \begin{cases} B(p) & \text{if } X = 1 \\ B(q) & \text{if } X = 0 \end{cases}$$

► $f_1 : \{0, 1\} \rightarrow \{0, 1\} :$

$$f = \begin{cases} 1 & \text{if } x = 1 \\ 0 & \text{if } x = 0 \end{cases}$$

$$\begin{aligned} R(f_1) &= E[I(Y, f(X))] \\ &= 1 \times P(Y \neq f(X)) + 0 \times P(Y = f(X)) \\ &= P(Y \neq f(X)) \\ &= P((Y \neq f(X)) \cap (X = 1)) + P((Y \neq f(X)) \cap (X = 0)) \end{aligned} \tag{4}$$

$$\begin{aligned} R(f_1) &= E[I(Y, f(X))] \\ &= 1 \times P(Y \neq f(X)) + 0 \times P(Y = f(X)) \\ &= P(Y \neq f(X)) \\ &= P((Y \neq f(X)) \cap (X = 1)) + P((Y \neq f(X)) \cap (X = 0)) \\ &= P((Y \neq f(X)) | X = 1)P(X = 1) \\ &\quad + P((Y \neq f(X)) | X = 0)P(X = 0) \end{aligned} \tag{5}$$

$$\begin{aligned} R(f_1) &= E[I(Y, f(X))] \\ &= 1 \times P(Y \neq f(X)) + 0 \times P(Y = f(X)) \\ &= P(Y \neq f(X)) \\ &= P((Y \neq f(X)) \cap (X = 1)) + P((Y \neq f(X)) \cap (X = 0)) \\ &= P((Y \neq f(X)) | X = 1)P(X = 1) \\ &\quad + P((Y \neq f(X)) | X = 0)P(X = 0) \\ &= \frac{1}{2}P((Y \neq 1) | X = 1) + \frac{1}{2}P((Y \neq 0) | X = 0) \end{aligned} \tag{6}$$

$$\begin{aligned}R(f_1) &= E[I(Y, f(X))]\nonumber\\&= 1 \times P(Y \neq f(X)) + 0 \times P(Y = f(X))\nonumber\\&= P(Y \neq f(X))\nonumber\\&= P((Y \neq f(X)) \cap (X = 1)) + P((Y \neq f(X)) \cap (X = 0))\nonumber\\&= P((Y \neq f(X))|X = 1)P(X = 1)\nonumber\\&\quad + P((Y \neq f(X))|X = 0)P(X = 0)\nonumber\\&= \frac{1}{2}P((Y = 0)|X = 1) + \frac{1}{2}P((Y = 1)|X = 0)\nonumber\\&= \frac{1}{2}(1 - p) + \frac{1}{2}q\end{aligned}$$

(7)

Exercise 2: Now consider

$$f_2 = \begin{cases} 0 & \text{if } x = 1 \\ 1 & \text{if } x = 0 \end{cases}$$

What is $R(f_2)$?

Exercise 2:

$$\forall x, f_2(x) = 1 - f_1(x) \quad (8)$$

Exercise 2 :

$$\forall x, f_2(x) = 1 - f_1(x) \quad (9)$$

Hence

$$\begin{aligned} R(f_2) &= P(Y \neq f_2(X)) \\ &= P(Y \neq (1 - f_1(X))) \\ &= P(Y = f_1(X)) \\ &= 1 - R(f_1) \end{aligned} \quad (10)$$

Exercise 3: Third predictor :

$$\forall x, f_3(x) = 1 \quad (11)$$

What is $R(f_3)$?

Exercice 3 :

$$\begin{aligned} R(f_3) &= P(Y \neq f_3(X)) \\ &= P(Y = 0) \end{aligned} \tag{12}$$

Exercise 3:

$$\begin{aligned} R(f_3) &= P(Y \neq f_3(X)) \\ &= P(Y = 0) \\ &= P(Y = 0 \cap X = 0) + P(Y = 0 \cap X = 1) \\ &= P(Y = 0|X = 0)P(X = 0) + P(Y = 0|X = 1)P(X = 1) \\ &= \frac{1}{2}(1 - p) + \frac{1}{2}(1 - q) \end{aligned} \tag{13}$$

Exercise 4 :

Now, we observe the following dataset :

$$D_4 = \{(0, 1), (0, 0), (0, 0), (1, 0)\} \quad (14)$$

Compute the empirical risks $R_4(f_1)$, $R_4(f_2)$, $R_4(f_3)$.

$$R_n(f) = \frac{1}{n} \sum_{i=1}^n l(y_i, f(x_i))$$

$$D_4 = \{(0, 1), (0, 0), (0, 0), (1, 0)\} \quad (15)$$

$$\begin{aligned} R_4(f_1) &= \frac{1}{4} \sum_{i=1}^4 l(f_1(x_i), y_i) \\ &= \frac{1}{4} \left(l(f_1(0), 1) + l(f_1(0), 0) + l(f_1(0), 0) + l(f_1(1), 0) \right) \\ &= \frac{1}{4} \times 2 \\ &= \frac{1}{2} \end{aligned} \quad (16)$$

$$D_4 = \{(0, 1), (0, 0), (0, 0), (1, 0)\} \quad (17)$$

$$\begin{aligned} R_4(f_2) &= \frac{1}{4} \sum_{i=1}^4 l(f_2(x_i), y_i) \\ &= \frac{1}{4} \left(l(f_2(0), 1) + l(f_2(0), 0) + l(f_2(0), 0) + l(f_2(1), 0) \right) \\ &= \frac{1}{4} \times 2 \\ &= \frac{1}{2} \end{aligned} \quad (18)$$

$$D_4 = \{(0, 1), (0, 0), (0, 0), (1, 0)\} \quad (19)$$

$$\begin{aligned} R_4(f_3) &= \frac{1}{4} \sum_{i=1}^4 l(f_3(x_i), y_i) \\ &= \frac{1}{4} \left(l(f_3(0), 1) + l(f_3(0), 0) + l(f_3(0), 0) + l(f_3(1), 0) \right) \\ &= \frac{1}{4} \times 3 \\ &= \frac{3}{4} \end{aligned} \quad (20)$$

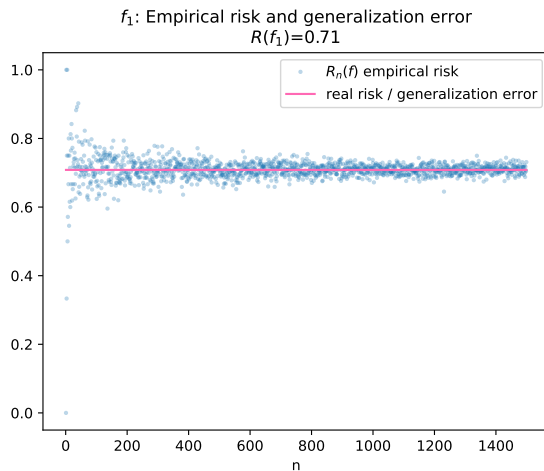
Random variable

- ▶ $R_4(f)$ (empirical risk) **depends** on D_4 . If we sample another dataset, $R_4(f)$ is likely to change, it is a **random variable**.
- ▶ $R(f)$ (generalization error) is **deterministic**, given the joint law of (X, Y) .


- └ Risks and risk decompositions
 - └ Expected value of empirical risk

Given a predictor f , a natural question arises :
Does $R_n(f)$ have a limit when $n \rightarrow +\infty$?


Simulations



Simulations



Simulations



Convergence of empirical risk

We fix $f \in H$ (hypothesis space). We assume that the samples (X_i, Y_i) are i.i.d, with the distribution of (X, Y) , noted ρ . Then, under some assumptions (for instance, if the empirical risks are bounded), we have that in probability :

$$\lim_{n \rightarrow +\infty} R_n(f) = R(f) \quad (21)$$

The empirical risk of a fixed f converges to its real risk.

Proof



$$R_n(f) = \frac{1}{n} \sum_{i=1}^n l(y_i, f(x_i))$$



$$\forall i, E[l(f(X_i), Y_i)] = E[l(f(X), Y)] \quad (22)$$

- ▶ i.i.d. variables.
- ▶ Law of large numbers.

Also

$$\begin{aligned} E_{D_n \sim \rho}(R_n(h)) &= \frac{1}{n} \sum_{i=1}^n E_{D_n \sim \rho}(l(f(X_i), Y_i)) \\ &= \frac{1}{n} \sum_{i=1}^n E_{(X, Y) \sim \rho}(l(f(X), Y)) \\ &= E_{(X, Y) \sim \rho}(l(f(X), Y)) \\ &= R(h) \end{aligned}$$

However, we do **not** have

$$E[R_n(\tilde{f}_n)] = R(\tilde{f}_n) \quad (23)$$

where \tilde{f}_n is the minimizer of the empirical risk.
 \tilde{f}_n depends on the dataset D_n .

$$E_{D_n \sim \rho}(r(\tilde{f}_n(X_i), Y_i)) \neq E_{(X, Y) \sim \rho}(r(\tilde{f}_n(X), Y)) \quad (24)$$

Risk decomposition

- ▶ f^* : Bayes predictor
- ▶ F : Hypothesis space
- ▶ \tilde{f}_n : estimated predictor ($\in F$).

$$E[R(\tilde{f}_n)] - R^* = \left(E[R(\tilde{f}_n)] - \inf_{f \in F} R(f) \right) + \left(\inf_{f \in F} R(f) - R^* \right) \quad (25)$$

Underfitting and overfitting

Approximation error (bias term) : depends on f^* and F , not on \tilde{f}_n , D_n .

$$\inf_{f \in F} R(f) - R^* \geq 0$$

Estimation error (variance term, fluctuation error, stochastic error) : depends on D_n , F , \tilde{f}_n .

$$E(R(\tilde{f}_n)) - \inf_{f \in F} R(f) \geq 0$$

- ▶ too small F : underfitting (large bias, small variance)
- ▶ too large F : overfitting (small bias, large variance)

Deterministic bound on the estimation error

We consider the best estimator in hypothesis space

$$f_a = \arg \min_{h \in F} R(h)$$

We can show that

$$R(\tilde{f}_n) - R(f_a) \leq 2 \sup_{h \in F} |R(h) - R_n(h)| \quad (26)$$

Deterministic bound on the estimation error

$$f_a = \arg \min_{h \in F} R(h)$$

$$\begin{aligned} R(\tilde{f}_n) - R(f_a) &= (R(\tilde{f}_n) - R_n(\tilde{f}_n)) \\ &\quad + (R_n(\tilde{f}_n) - R_n(f_a)) \\ &\quad + (R_n(f_a) - R(f_a)) \end{aligned} \tag{27}$$

Deterministic bound on the estimation error

$$f_a = \arg \min_{h \in F} R(h)$$

$$\begin{aligned} R(\tilde{f}_n) - R(f_a) &= (R(\tilde{f}_n) - R_n(\tilde{f}_n)) \\ &\quad + (R_n(\tilde{f}_n) - R_n(f_a)) \\ &\quad + (R_n(f_a) - R(f_a)) \end{aligned} \tag{28}$$

But by definition \tilde{f}_n minimizes R_n , so $(R_n(\tilde{f}_n) - R_n(f_a)) \leq 0$.

Deterministic bound on the estimation error

$$R(\tilde{f}_n) - R(f_a) \leq 2 \sup_{h \in F} |R(h) - R_n(h)| \quad (29)$$

Later in the course, based on **concentration inequalities** we will further build on this result and prove a probabilistic bound of the form

$$R(\tilde{f}_n) - R(f_a) \leq \frac{C}{\sqrt{n}} \quad (30)$$

(remember that by definition $0 \leq R(\tilde{f}_n) - R(f_a)$)

Order of magnitude of estimation error

We keep in mind that

$$R(\tilde{f}_n) - R(f_a) = \mathcal{O}\left(\frac{C}{\sqrt{n}}\right) \quad (31)$$

Approximate solution

- ▶ In machine learning, it is often not necessary to find the actual minimizer of the empirical risk , as there is an estimation error of $\mathcal{O}(\frac{1}{\sqrt{n}})$. [Bottou and Bousquet, 2009,]
- ▶ We can use an approximate solution \hat{f}_n , such that

$$R_n(\hat{f}_n) \leq R_n(\tilde{f}_n) + \rho \quad (32)$$

with ρ a predefined tolerance.

- ▶ This important because in large-scale ML, the **computation time** need to be optimized.

Approximate solution

This gives a new risk decomposition :

$$\begin{aligned} E[R(\hat{f}_n)] - R^* &= \left(E[R(\hat{f}_n)] - E[R(\tilde{f}_n)] \right) \\ &\quad + \left(E[R(\tilde{f}_n)] - \inf_{f \in F} R(f) \right) \\ &\quad + \left(\inf_{f \in F} R(f) - R^* \right) \end{aligned} \tag{33}$$

Approximate solution

This gives a new risk decomposition :

$$\begin{aligned} E[R(\hat{f}_n)] - R^* &= \left(E[R(\hat{f}_n)] - E[R(\tilde{f}_n)] \right) \\ &\quad + \left(E[R(\tilde{f}_n)] - \inf_{f \in F} R(f) \right) \\ &\quad + \left(\inf_{f \in F} R(f) - R^* \right) \end{aligned} \quad (34)$$

$E[R(\hat{f}_n)] - E[R(\tilde{f}_n)]$ is the **optimization error**.

To conclude, we have :

- ▶ an approximation error
- ▶ an estimation error
- ▶ an optimization error

Risks and risk decompositions

Examples

Expected value of empirical risk

Risk decomposition

Optimization error

Optimization in machine learning

Existence results

Convex analysis

Gradients

Ordinary Least squares II

OLS estimator

Statistical analysis of OLS

Minimizers

Definition

Let $f : \mathbb{R}^d \rightarrow \mathbb{R}$ be defined on $K \subset \mathbb{R}^d$.

$x \in K$ is a local minimum of f on K if and only if

$$\exists \delta > 0, \forall y \in K, \|y - x\| < \delta \Rightarrow f(x) \leq f(y)$$

$x \in K$ is a global minimum of f on K if and only if

$$\forall y \in K, f(x) \leq f(y)$$

Existence result

Theorem

Existence of a global minimum in \mathbb{R}^d

Let K be a closed non-empty subset of \mathbb{R}^d , and $f : \mathbb{R}^d \rightarrow \mathbb{R}$ a continuous coercive function. Then, there exists at least a global minimum of f on K .

Convexity

Definition

The function $f : \Omega \rightarrow \mathbb{R}$ with Ω convex is :

- **convex** if $\forall x, y \in \Omega, \alpha \in [0, 1]$,

$$f(\alpha x + (1 - \alpha)y) \leq \alpha f(x) + (1 - \alpha)f(y)$$

- **strictly convex** if $\forall x, y \in \Omega, \alpha \in [0, 1]$,

$$f(\alpha x + (1 - \alpha)y) < \alpha f(x) + (1 - \alpha)f(y)$$

- **μ -strongly convex** if $\forall x, y \in \Omega, \alpha \in [0, 1]$,

$$f(\alpha x + (1 - \alpha)y) \leq \alpha f(x) + (1 - \alpha)f(y) - \frac{\mu}{2}\alpha(1 - \alpha)\|x - y\|^2$$

Examples

- ▶ All norms are convex.
- ▶ $x \mapsto \theta^T x$ is convex on \mathbb{R}^d with $\theta \in \mathbb{R}^d$ (linear form)
- ▶ if Q is a symmetric semidefinite positive matrix, then $x \mapsto x^T Q x$ is convex.
- ▶ if Q is a symmetric definite positive matrix (matrice définie positive) with smallest eigenvalue $\lambda_{\min} > 0$, then $x \mapsto x^T Q x$ is $2\lambda_{\min}$ -strongly convex.
- ▶ If f is increasing and convex and g is convex, then $f \circ g$ is convex.
- ▶ If f is convex and g is linear, then $f \circ g$ is convex.

Differential formulation of convexity

Proposition

Let $f : V \rightarrow \mathbb{R}$ be a differentiable function. The following conditions are equivalent.

- ▶ *f is convex.*
- ▶ $\forall x, y \in V, f(y) \geq f(x) + (f'(x)|y - x)$ (*f is above its tangent space*)
- ▶ $\forall x, y \in V, (f'(x) - f'(y)|x - y) \geq 0$ (*f' grows*)

Differential formulation of strong convexity

Proposition

Let $f : V \rightarrow \mathbb{R}$ be a differentiable function, and $\mu > 0$. The following conditions are equivalent.

- ▶ f is μ -convex
- ▶ $\forall x, y \in V, f(y) \geq f(x) + (f'(x)|y - x) + \frac{\mu}{2}\|y - x\|^2$
- ▶ $\forall x, y \in V, (f'(x) - f'(y)|x - y) \geq \mu\|x - y\|^2$

Convexity of two-times differentiable functions

- ▶ f is convex if and only if

$$\forall x, h \in y, J''(x)(h, h) \geq 0$$

- ▶ f is μ -strongly convex if and only if

$$\forall x, h \in y, J''(x)(h, h) \geq \mu \|h\|^2$$

Convexity and Hessian

If $V = \mathbb{R}^d$, this translates into

$$\forall x, h \in y, h^T (H_x f) h \geq 0 \quad (35)$$

and


$$\forall x, h \in y, h^T (H_x f) h \geq \mu \|h\|^2 \quad (36)$$

- ▶ 35 means that $\forall x \in \mathbb{R}^d$, all eigenvalues of $H_x f$ are non-negative (positive semi-definite Hessian)
- ▶ 36 means that they all are $\geq \mu$ (positive definite Hessian).

Positive semi-definite Hessian

$$H_x^f = \begin{pmatrix} 2 & 0 \\ 0 & 0 \end{pmatrix} \quad (37)$$


Positive semi-definite Hessian



Positive definite Hessian

$$H_x^f = \begin{pmatrix} 2 & 0 \\ 0 & 2 \end{pmatrix} \quad (38)$$

Positive definite Hessian



Minima of convex functions

Proposition

- ▶ *If f is convex, any local minimum is a global minimum. The set of global minimizers is a convex set.*
- ▶ *If f is strictly convex, there exists at most one local minimum (that is thus global).*
- ▶ *If f is convex and C^1 (differentiable, $a \mapsto df_a$ continuous), then x is a minimum (thus global) of f on V if and only if the gradient cancels in x , $\nabla_x f = 0$. V need not be finite-dimensional.*

Risks and risk decompositions

- Examples

- Expected value of empirical risk

- Risk decomposition

- Optimization error

Optimization in machine learning

- Existence results

- Convex analysis

- Gradients

Ordinary Least squares II

- OLS estimator

- Statistical analysis of OLS

OLS

- ▶ $\mathcal{X} = \mathbb{R}^d$
- ▶ $\mathcal{Y} = \mathbb{R}$.
- ▶ $l(y, y') = (y - y')^2$ (squares loss)
- ▶

$$F = \{x \mapsto \theta^T x, \theta \in \mathbb{R}^d\}$$

OLS

The dataset is stored in the **design matrix** $X \in \mathbb{R}^{n \times d}$.

$$X = \begin{pmatrix} x_1^T \\ \dots \\ x_i^T \\ \dots \\ x_n^T \end{pmatrix} = \begin{pmatrix} x_{11}, \dots, x_{1j}, \dots, x_{1d} \\ \dots \\ x_{i1}, \dots, x_{ij}, \dots, x_{id} \\ \dots \\ x_{n1}, \dots, x_{nj}, \dots, x_{nd} \end{pmatrix}$$

The vector of predictions of the estimator writes $Y = X\theta$. Hence,

$$\begin{aligned} R_n(\theta) &= \frac{1}{n} \sum_{i=1}^n (y_i - \theta^T x_i)^2 \\ &= \frac{1}{n} \|Y - X\theta\|_2^2 \end{aligned}$$

OLS estimator

Exercise 5 : Convexity

Show that the objective function $R_n(\theta)$ is convex in θ .

$$\theta \mapsto \|Y - X\theta\|_2^2 \quad (39)$$

OLS estimator

We assume that X is **injective**. Necessary, $d \leq n$.

Proposition

Closed form solution

*We X is injective, there exists a unique minimiser of $R_n(\theta)$, called the **OLS estimator**, given by*

$$\hat{\theta} = (X^T X)^{-1} X^T Y \quad (40)$$

Setup

Assumptions :

- ▶ **Linear model** : $\exists \theta^* \in \mathbb{R}^d$,

$$Y_i = \theta^{*T} x_i + Z_i, \forall i \in [1, n]$$

and Z_i is a centered noise (or error) ($E[Z_i] = 0$) with variance σ^2 .

- ▶ Fixed design X .

In this setup, we wonder :

- ▶ 1) what is the Bayes predictor ? What is the Bayes risk ?
- ▶ 2) is the expected value of OLS equal to the Bayes predictor ?
- ▶ 3) what is the excess risk of the OLS estimator ?

1) Bayes predictor

With the square loss, we always have that the Bayes predictor is the conditional expectation, see FTML.pdf section 3.1.3.

$$f^*(x) = E[Y|X = x] \quad (41)$$

1) Bayes predictor

$$\begin{aligned}f^*(x) &= E[Y|X = x] \\&= E[X^T \theta^* + \epsilon | X = x] \\&= E[X^T \theta^* | X = x] + E[\epsilon | X = x] \\&= X^T \theta^*\end{aligned}\tag{42}$$

1) Bayes risk

$$\begin{aligned} R^* &= E_{X,Y}[(Y - f^*(X))^2] \\ &= E_{X,\epsilon}[(X^T \theta^* + \epsilon - X^T \theta^*)^2] \\ &= E_{X,\epsilon}[\epsilon^2] \\ &= \sigma^2 \end{aligned} \tag{43}$$


2) Expected value of $\hat{\theta}$

$$\begin{aligned}E[\hat{\theta}] &= E[(X^T X)^{-1} X^T Y] \\&= E[(X^T X)^{-1} X^T (X\theta^* + \epsilon)] \\&= E[(X^T X)^{-1} X^T (X\theta^*)] + E[(X^T X)^{-1} X^T \epsilon] \\&= E[(X^T X)^{-1} (X^T X) \theta^*] + (X^T X)^{-1} X^T E[\epsilon] \\&= E[\theta^*] \\&= \theta^*\end{aligned}\tag{44}$$

We conclude that the OLS estimator is an **unbiased estimator** of θ^* .

3) Excess risk + variance


$$E[R(\hat{\theta})] - R(\theta^*) = \frac{\sigma^2 d}{n} \quad (45)$$



4) Variance

$$\text{Var}(\hat{\theta}) = \frac{\sigma^2}{n} \Sigma^{-1} \quad (46)$$

with $\Sigma = X^T X \in \mathbb{R}^{d \times d}$.



References I



Bottou, L. and Bousquet, O. (2009).

The tradeoffs of large scale learning.

*Advances in Neural Information Processing Systems 20 -
Proceedings of the 2007 Conference*, (January 2007).