

FTML practical session 2: 2024/03/15




TABLE DES MATIÈRES

1	Linear regression	2
1.1	A simple example	2
1.2	Formalization	2
1.3	Empirical risk	3
2	Ordinary least squares : empirical risk and overfitting	3
2.1	OLS estimator	3
2.2	Excess risk the OLS estimator	4
2.3	Statistical setting	4
2.4	Impact of the dimensions on the OLS overfitting	4

INTRODUCTION

The goal of this session is to keep exploring the concepts linked to risks : Bayes risk, empirical risk, overfitting. We keep exploring them in a linear regression context, but dive a bit deeper technically and start studying the influence of the dimensions (number of samples, number of features) on the statistical properties of the results. Some aspects of this session are a bit technical, it is normal if you need to take some time to all understand the objects used in the session.

There are some template files that you can use in the repo, but you can also start from scratch with your own files if you prefer.

- `main_ols.py`
- `utils_plots.py`
- `utils_algo.py`
- `constants.py`

1 LINEAR REGRESSION

A **linear model**, such as the Ordinary least squares (OLS), can be interpreted as predicting an output value (dependent variable) from combining the contributions from the d **features** of the input data (independent variables), in a linear way. This can be useful for classification as well as regression.

1.1 A simple example

For instance, if I want to predict the amount of money that I will spend when buying some clothes, I can use a linear model. If θ contains the price of each type of clothe, and x the number of each type of clothe that I buy, then I have to spend $x^T \theta$. If there exists $d = 4$ types of clothes with a price θ_i :

- socks : $\theta_1 = 2$
- T-shirts : $\theta_2 = 25$
- pants : $\theta_3 = 50$
- hats : $\theta_4 = 20$

If I want to buy 10 socks, 2 T-shirts, 1 pants and 1 hat, then $x^T = (10, 2, 1, 1)$ and I spend

$$\begin{aligned} x^T \theta &= 10 \times 2 + 2 \times 25 + 1 \times 50 + 1 \times 20 \\ &= 140 \end{aligned} \tag{1}$$

Obviously, not all phenomena can be approximated well in a linear way. However, linear regression is a foundation for more advanced modelisation that we will study in future classes (feature maps, kernel methods, neural networks, etc).

As we will study the influence of the dimensions of the problem n (number of samples) and d (number of features of each sample), we will work with abstract datasets, instead of a given "physical dataset". However, you can always refer to the previous example in order to interpret linear models that we will use.

1.2 Formalization

Let us abstract the notations a little bit. In the OLS setting,

- $\mathcal{X} = \mathbb{R}^d$ (input space)
- $\mathcal{Y} = \mathbb{R}$ (output space)

An input sample is stored in a column vector $x_a \in \mathbb{R}^d$, with $a \in \mathbb{N}$ like

$$x_a = \begin{pmatrix} x_{a1} \\ \dots \\ x_{ai} \\ \dots \\ x_{ad} \end{pmatrix} \in \mathbb{R}^d \quad (2)$$

The estimator is a **linear mapping** parametrized by $\theta \in \mathbb{R}^d$. The prediction associated with $x \in \mathbb{R}^d$ is $f(x) = \theta^T x = x^T \theta$. If we use the squared-loss, the discrepancy between two real numbers y and y' writes $l(y, y') = (y - y')^2$. Finally, the input dataset is stored in the **design matrix** $X \in \mathbb{R}^{n \times d}$.

$$X = \begin{pmatrix} x_1^T \\ \dots \\ x_i^T \\ \dots \\ x_n^T \end{pmatrix} = \begin{pmatrix} x_{11}, \dots, x_{1j}, \dots, x_{1d} \\ \dots \\ x_{i1}, \dots, x_{ij}, \dots, x_{id} \\ \dots \\ x_{n1}, \dots, x_{nj}, \dots, x_{nd} \end{pmatrix} \quad (3)$$

and the output labels are stored in a vector

$$y = \begin{pmatrix} y_1 \\ \dots \\ y_i \\ \dots \\ y_n \end{pmatrix} \in \mathbb{R}^n \quad (4)$$

1.3 Empirical risk

The **empirical risk** of an estimator θ (when we talk about an estimator, it is here equivalent to refer to θ directly or to the mapping defined by θ)

$$R_n(\theta) = \frac{1}{n} \sum_{i=1}^n (y_i - \theta^T x_i)^2 \quad (5)$$

How can we write the empirical risk using only matrices, vectors and an euclidean norm?

By contrast, the **generalization error** of an estimator does not depend on the dataset. It is a fixed number, defined by (in this context of linear regression, squared loss)

$$R(\theta) = E[(y - \theta^T x)^2] \quad (6)$$

where the expected value is taken over the law of x and y , the random variables representing the input and output respectively. This law is unknown, in general.

2 ORDINARY LEAST SQUARES : EMPIRICAL RISK AND OVERFITTING

2.1 OLS estimator

The **OLS estimator**, noted $\hat{\theta}$, is the value of θ that minimizes $R_n(\theta)$. It is thus the solution to the problem of **empirical risk minimization**, a standard approach in supervised learning. We admit the following proposition :

Proposition. *Closed form solution*

If X is injective, there exists a unique minimiser of $R_n(\theta)$, called the **OLS estimator**, given by

$$\hat{\theta} = (X^T X)^{-1} X^T y \quad (7)$$

2.2 Excess risk the OLS estimator

Given an estimator, we are interested in its **excess risk**. It is the difference between its generalization error and the bayes risk (which is the definition of overfitting that we used this morning replacing the least squares error by the R_2 scores). We will evaluate the excess risk of the OLS estimator as a function of n and d , for a given statistical setting, in order to observe the results of figure 1.

2.3 Statistical setting

In order to compute these quantities, it is necessary to make statistical assumptions. We will use the **linear model**, with **fixed design**, a classical framework to analyze OLS. In this setting, we assume that X is given and fixed, and that there exists a vector $\theta^* \in \mathbb{R}^d$, such that $\forall i \in \{1, \dots, n\}$,

$$y_i = x_i^T \theta^* + \epsilon_i \quad (8)$$

where for all $i \in \{1, \dots, n\}$, ϵ_i are independent, with expectation $E[\epsilon_i] = 0$ and variance $E[\epsilon_i^2] = \sigma^2$. The ϵ_i represent a variability in the output, that is due to **noise**, or to the presence of unobserved variables. Put together in a vector $\epsilon \in \mathbb{R}^n$, this allows to write

$$y = X\theta^* + \epsilon \quad (9)$$

Considering this morning's practical session observations, what do we expect to be the Bayes estimator?

We admit (but you can also prove it) that the corresponding Bayes risk is σ^2 .

For a given n and d (with $d \leq n$), run a simulation that reproduces the fixed design, linear model statistical setting and verify that the estimator that you have chosen achieves the Bayes risk. Verify that different estimators have worse (larger) risks. You need to generate first a X and a θ^* .

You will again need to approximate the real risk by the empirical risk, a number of times that is sufficient to observe the convergence of the law of large numbers.

2.4 Impact of the dimensions on the OLS overfitting

We now know that in this statistical setting, the Bayes risk is σ^2 . In practical applications, we of course do not have access to θ^* and we must find a relevant θ , based on the data only. We must resort to empirical risk minimization.

We want to compare the risk of the OLS estimator to the Bayes risk, to know how good the OLS estimator compares to the Bayes estimator. Again, remember that in practical applications, we will not have access to the Bayes estimator or the Bayes risk.

Run a simulation that generates an output vector y , computes the corresponding OLS estimator $\hat{\theta}$ and an estimation of its generalization error by sampling more outputs and by computing the empirical risk of $\hat{\theta}$ on these new outputs.

In order to generate the X matrices in various dimensions, you can use uniformly distributed entries. This should ensure that X is injective.

Run the same simulation for various values of n and d in order to reproduce the results of figure 1.

We will prove during the lectures the different results mentioned in this exercise.






FIGURE 1 – Dependence of the risk (generalization error)

If $d > n$, the $\sigma^2 d/n$ law is not respected!

Why can we still output an OLS estimator when $d > n$?

FIGURE 2 – Test errors with cases where $d > n$.

Optional : monitor the stability of the OLS estimator as a function of n and d , for instance by computing the sample standard deviation of the random variable $\|\hat{\theta} - \theta^*\|/\|\theta^*\|$ as a function of n and d . When d is closer to n , $\hat{\theta}$ can become very unstable (a small variation of the noise might lead to a large variation of $\hat{\theta}$). This behavior depends on the conditioning of $X^T X$, as we will see later.)

FIGURE 3 – Instability of the OLS estimator when d gets close to n .