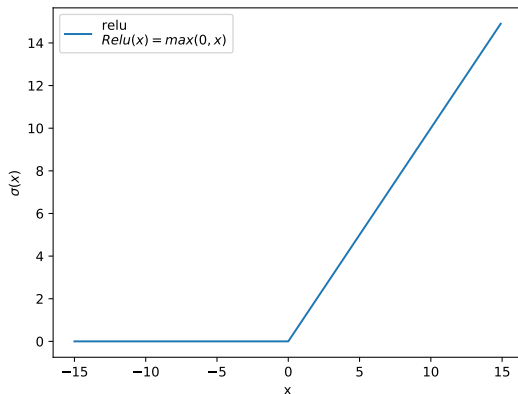


# Fondamentaux théoriques du machine learning



# Overview of lecture 7

## Optimization of neural networks

- Difficulties of optimizing neural networks
- Specific methods for neural networks

## Statistical learning

- Bounding the estimation error
- Interpolation regime and double descent

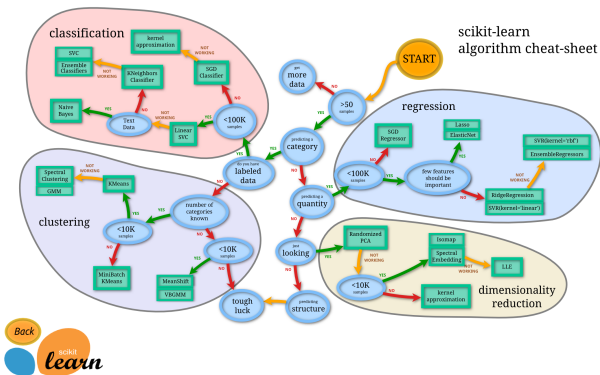
## Local averaging methods

- Supervised learning
- Density estimation

## Adaptivity

- No free lunch theorems
- Adaptivity

## ML map



[https://scikit-learn.org/stable/tutorial/machine\\_learning\\_map/index.html](https://scikit-learn.org/stable/tutorial/machine_learning_map/index.html)

## Optimization of neural networks

- Difficulties of optimizing neural networks

- Specific methods for neural networks

## Statistical learning

- Bounding the estimation error

- Interpolation regime and double descent

## Local averaging methods

- Supervised learning

- Density estimation

## Adaptivity

- No free lunch theorems

- Adaptivity

# Neural networks

References / tools :

- ▶ <https://www.deeplearningbook.org/>
- ▶ <https://d2l.ai/>
- ▶ [https://mlelarge.github.io/dataflowr-web/dldiy\\_ens.html](https://mlelarge.github.io/dataflowr-web/dldiy_ens.html)
- ▶ <https://playground.tensorflow.org/>
- ▶ <http://www.jzliu.net/blog/simple-python-library-visualize-neural-network/>

## Learning representations / features

►  $\mathcal{X} = \mathbb{R}^d$ .

►  $\mathcal{Y} = \mathbb{R}$ .

A neural network with scalar output learns a feature map  $\phi : \mathbb{R}^d \rightarrow \mathbb{R}^m$  **and** a linear regressor or classifier,  $\theta \in \mathbb{R}^m$ .

$$\forall x, f(x) = \langle \theta, \phi(x) \rangle \quad (1)$$

We can add a bias by adding a dimension to  $\theta$  and adding a component with a 1 to each  $\phi(x)$ .

## Learning representations / features

►  $\mathcal{X} = \mathbb{R}^d$ .

►  $\mathcal{Y} = \mathbb{R}$ .

A neural network learns a feature map  $\phi : \mathbb{R}^d \rightarrow \mathbb{R}^m$  **and** a linear regressor or classifier,  $\theta \in \mathbb{R}^m$ .

$$\forall x, f(x) = \langle \theta, \phi(x) \rangle \quad (2)$$

We can add a bias by adding a dimension to  $\theta$  and adding a component with a 1 to each  $\phi(x)$ .

**Remark** : kernel methods use hardcoded features  $\phi$ , that can be **implicit** (kernel trick).

## Multi output

In order to learn a multidimensional output of dimension  $p$  ( $\mathcal{Y} = \mathbb{R}^p$ ), it is sufficient to learn a matrix  $\theta \in \mathbb{R}^{m,p}$ .

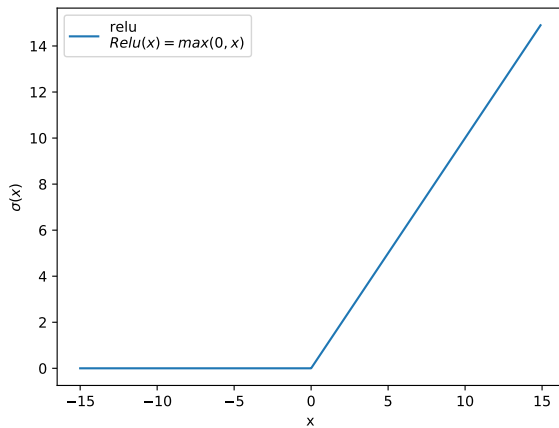


## Single layer neural network

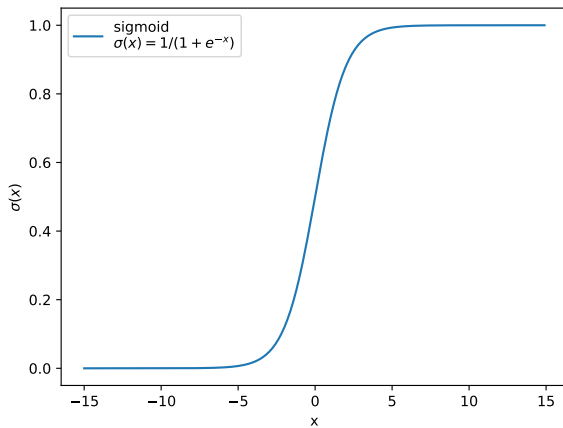
$$f(x) = \sum_{j=1}^m \theta_j \sigma(w_j^T x + b_j) \quad (3)$$

$\sigma$  is an activation function (sigmoid, tanh, ReLu, etc).

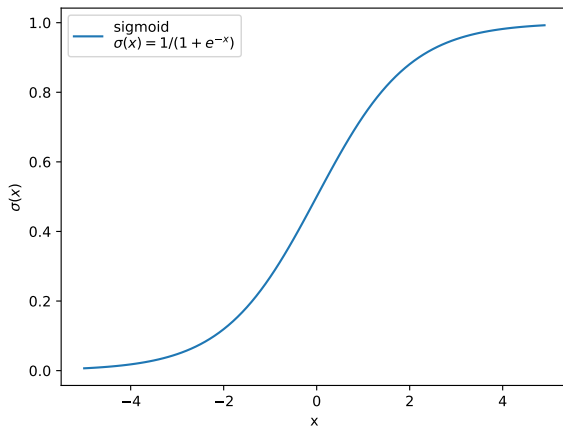
# ReLU



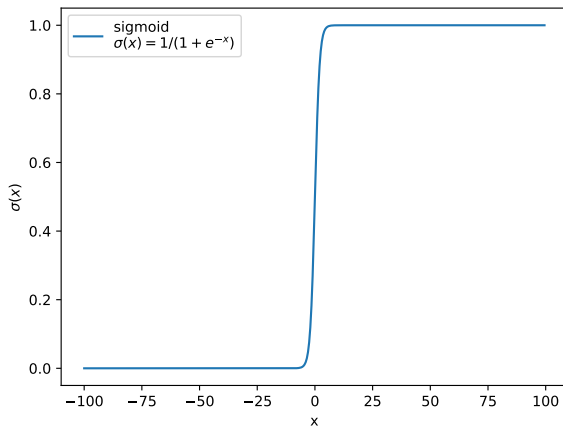
# Sigmoid



# Sigmoid



# Sigmoid



# Automatic differentiation

When working with neural networks, most used libraries implement automatic differentiation.

- ▶ tensorflow
- ▶ pytorch (autograd)

## Optimizing neural networks

Optimizing neural networks comes with specific difficulties.

- ▶ the problem is non-convex
- ▶ there is often a large number of parameters (optimization in a high dimensional space)
- ▶ specific problems due to depth (vanishing gradients, see below)

<http://yann.lecun.com/exdb/publis/pdf/lecun-98b.pdf>

## Non-convexity

We know that

- ▶ If  $f$  is increasing and convex and  $g$  is convex, then  $f \circ g$  is convex.
- ▶ If  $f$  is convex and  $g$  is linear, then  $f \circ g$  is convex.

With neural networks, we are in neither of these cases, as the activations  $\sigma$  are non linear.



## Non-convexity

We know that


- ▶ If  $f$  is increasing and convex and  $g$  is convex, then  $f \circ g$  is convex.
- ▶ If  $f$  is convex and  $g$  is linear, then  $f \circ g$  is convex.

With neural networks, we are in neither of these cases, as the activations  $\sigma$  are non linear.

Hence **the objective function is non-convex**, and it remains difficult to understand why gradient based methods often perform well in practice

- └ Optimization of neural networks
- └ Difficulties of optimizing neural networks

# Vanishing gradients



## Exercise 1 :

What is the maximum value of  $|\sigma'(z)|$  ?

## Exponentially decreasing gradients

- ▶ At each layer, the gradients are multiplied by a term of the form  $\sigma'(u)$ . Using a large number of layers leads to gradient norms that decrease rapidly when we move away from the output layer.
- ▶ This slows training down and caused deep learning to plateau for some years.
- ▶ Several initializations were necessary in order to obtain convergence, the result was unstable.

- └ Optimization of neural networks
  - └ Difficulties of optimizing neural networks

# ReLU

The usage of ReLU solved this problem.

Other activation functions :

[https://dashee87.github.io/deep%20learning/  
visualising-activation-functions-in-neural-networks/](https://dashee87.github.io/deep%20learning/visualising-activation-functions-in-neural-networks/)

## SGD variants for neural networks

Several specific variations of SGD are commonly used for deep learning.

<https://pytorch.org/docs/stable/optim.html>

# Specific methods for deep learning

Architectures :

- ▶ Convotutional networks
- ▶ Residual neural network (ResNet)

Optimization / regularization :

- ▶ dropout
- ▶ batch normalisation

## Optimization of neural networks

Difficulties of optimizing neural networks

Specific methods for neural networks

## Statistical learning

Bounding the estimation error

Interpolation regime and double descent

## Local averaging methods

Supervised learning

Density estimation

## Adaptivity

No free lunch theorems

Adaptivity

# Statistical learning

- ▶ We come back to the statistical analysis of supervised learning.
- ▶ More precisely, to that of empirical risk minimization.



## Reminder on risks

Let  $l$  be a loss. Generalization error :

$$R(f) = E_{(X,Y) \sim \rho}[l(Y, f(X))] \quad (4)$$

The **empirical risk (ER)** of an estimator  $f$  writes

$$R_n(f) = \frac{1}{n} \sum_{i=1}^n l(y_i, f(x_i)) \quad (5)$$

Remember that the risks depends on the loss  $l$ .

## Risk decomposition

- ▶  $f^*$  : Bayes predictor
- ▶  $F$  : Hypothesis space
- ▶  $f_n$  : estimated predictor (hence in  $F$ ).

$$E[R(f_n)] - R^* = \left( E[R(f_n)] - \inf_{f \in F} R(f) \right) + \left( \inf_{f \in F} R(f) - R^* \right) \quad (6)$$

## Risk decomposition

- ▶  $f^*$  : Bayes predictor
- ▶  $F$  : Hypothesis space
- ▶  $f_n$  : estimated predictor.

When doing empirical risk minimization,  $f_n$  is obtained by minimization of the empirical risk.

**However :**

- ▶ we have seen that in many cases, finding the exact minimizer of the empirical risk might be computationnaly hard.
- ▶ also, a natural question is whether it is sufficient to have an **approximate minimizer** of the empirical risk, as the empirical risk is an **approximation** of the generalization error.

## Risk decomposition

**Estimation error (variance term, fluctuation error, stochastic error)** : depends on  $D_n$ ,  $F$ ,  $f_n$ .

$$E\left[R(f_n)\right] - \inf_{f \in F} R(f) \geq 0$$

**Approximation error (bias term)** : depends on  $f^*$  and  $F$ , not on  $f_n$ ,  $D_n$ .

$$\inf_{f \in F} R(f) - R^* \geq 0$$

It is also possible to consider the **Optimization error** : depends on  $D_n$ ,  $F$ ,  $f_n$ .

$$E\left[R(\hat{f}_n) - R(f_n)\right] \tag{7}$$

where  $\hat{f}_n$  is an approximate solution to the optimization problem.

## Bound on the estimation error

We will now focus on the estimation error

$$E\left[R(f_n)\right] - \inf_{f \in F} R(f) \geq 0$$

$f_n$  is the empirical risk minimizer

We consider the best estimator in hypothesis space

$$f_a = \arg \min_{h \in F} R(h)$$

We have seen that

$$R(f_n) - R(f_a) \leq 2 \sup_{h \in F} |R(h) - R_n(h)| \quad (8)$$

# Deterministic bound on the estimation error

$$f_a = \arg \min_{h \in F} R(h) \quad (9)$$

$$f_n = \arg \min_{h \in F} R_n(h) \quad (10)$$

$$\begin{aligned} R(f_n) - R(f_a) &= (R(f_n) - R_n(f_n)) \\ &\quad + (R_n(f_n) - R_n(f_a)) \\ &\quad + (R_n(f_a) - R(f_a)) \end{aligned} \quad (11)$$

## Deterministic bound on the estimation error

$$f_a = \arg \min_{h \in F} R(h)$$

$$f_n = \arg \min_{h \in F} R_n(h) \quad (12)$$

$$\begin{aligned} R(f_n) - R(f_a) &= (R(f_n) - R_n(f_n)) \\ &\quad + (R_n(f_n) - R_n(f_a)) \\ &\quad + (R_n(f_a) - R(f_a)) \end{aligned} \quad (13)$$

But by definition  $f_n$  minimizes  $R_n$ , so  $(R_n(f_n) - R_n(f_a)) \leq 0$ .

## Deterministic bound on the estimation error

$$f_a = \arg \min_{h \in F} R(h)$$

$$f_n = \arg \min_{h \in F} R_n(h) \quad (14)$$

$$\begin{aligned} R(f_n) - R(f_a) &= (R(f_n) - R_n(f_n)) \\ &\quad + (R_n(f_n) - R_n(f_a)) \\ &\quad + (R_n(f_a) - R(f_a)) \end{aligned} \quad (15)$$

But by definition  $f_n$  minimizes  $R_n$ , so  $(R_n(f_n) - R_n(f_a)) \leq 0$ .

Finally :

$$R(f_n) - R(f_a) \leq 2 \sup_{h \in F} |R(h) - R_n(h)| \quad (16)$$



## Bound on the estimation error

If we are able to bound  $\sup_{h \in F} |R(h) - R_n(h)|$ , then we have a bound on the estimation error.

## Bound on the estimation error

### Theorem

*Weak law of large numbers*

*Let  $(X_i)_{i \in \mathbb{N}}$  be a sequence of i.i.d. variables that have a moment of order 2. We note  $m$  their expected value. Then*

$$\forall \epsilon > 0, \lim_{n \rightarrow +\infty} P\left(\left|\frac{1}{n} \sum_{i=1}^n X_i - m\right| \geq \epsilon\right) = 0 \quad (17)$$

*We say that we have **convergence in probability**.*

## Bound on the estimation error

### Theorem

*Weak law of large numbers*

*Let  $(X_i)_{i \in \mathbb{N}}$  be a sequence of i.i.d. variables that have a moment of order 2. We note  $m$  their expected value. Then*

$$\forall \epsilon > 0, \lim_{n \rightarrow +\infty} P\left(\left|\frac{1}{n} \sum_{i=1}^n X_i - m\right| \geq \epsilon\right) = 0 \quad (18)$$

*We say that we have **convergence in probability**.*

**However**, this is only an asymptotical result : it is a limit for  $n \rightarrow +\infty$ .

## Bound on the estimation error

If we are able to bound  $\sup_{h \in F} |R(h) - R_n(h)|$ , then we have a bound on the estimation error.

To do this, we will use some other mathematical results :

- ▶ Boole's inequality
- ▶ Hoeffding's inequality (non-asymptotical probabilistic bound)

## Boole's inequality

### Proposition

*Let  $A_1, A_2, \dots$ , be a countable set of events of a probability space  $\{\Omega, \mathcal{F}, P\}$ .*

*Then,*

$$P\left(\bigcup_{i \geq 1} A_i\right) \leq \sum_{i \geq 1} P(A_i) \quad (19)$$

This set might be infinite.

## Boole's inequality

### Proposition

*Let  $A_1, A_2, \dots$ , be a countable set of events of a probability space  $\{\Omega, \mathcal{F}, P\}$ .*

*Then,*

$$P\left(\bigcup_{i \geq 1} A_i\right) \leq \sum_{i \geq 1} P(A_i) \quad (20)$$

This set might be infinite.

**Exercise 2:** Prove the proposition.

# Hoeffding's inequality

## Theorem

### *Hoeffding's inequality*

Let  $(X_i)_{1 \leq i \leq n}$  be  $n$  i.i.d real random variables such that  $\forall i \in [1, n]$ ,  $X_i \in [a, b]$  and  $E(X_i) = \mu \in \mathbb{R}$ . Let  $\bar{\mu} = \frac{1}{n} \sum_{i=1}^n X_i$ .

Then  $\forall \epsilon > 0$ ,

$$P\left(|\bar{\mu} - \mu| \geq \epsilon\right) \leq 2 \exp\left(-\frac{2n\epsilon^2}{(b-a)^2}\right)$$

We admit this theorem.

# Setting

- ▶ Supervised learning.
- ▶ Finite space of estimator  $F$ .
- ▶ The loss  $l$  is uniformly bounded :  $l(\hat{y}, y) \in [a, b]$  with  $a$  and  $b$  real numbers.



# Step 1

We have seen that

$$R(f_n) - R(f_a) \leq 2 \sup_{h \in F} |R(h) - R_n(h)| \quad (21)$$

As a consequence, for all  $t \geq 0$  :

$$P\left(R(f_n) - R(f_a) \geq t\right) \leq P\left(2 \sup_{h \in F} |R(h) - R_n(h)| \geq t\right) \quad (22)$$

## Step 2

The fact that

$$2 \sup_{h \in F} |R(h) - R_n(h)| \geq t \quad (23)$$

is equivalent to :

$$\cup_{h \in F} \left( 2 |R(h) - R_n(h)| \geq t \right) \quad (24)$$

## Conclusion

**Exercise 3:** Using Boole's inequality and Hoeffding's inequality, show that

$$P\left(R(f_n) - R(f_a) \geq t\right) \leq 2|F| \exp\left(-\frac{nt^2}{2(b-a)^2}\right) \quad (25)$$

## Conclusion

We write

$$\delta = 2|F| \exp\left(-\frac{nt^2}{2(b-a)^2}\right) \quad (26)$$

Exercise 4 :

We assume that  $b - a = 1$ . Show that with probability  $\geq 1 - \delta$ ,

$$R(f_n) \leq R(f_a) + 2\sqrt{\frac{\log(|F|) + \log(\frac{2}{\delta})}{2n}} \quad (27)$$

# Generalization

It is possible to generalize to infinite sets :

- ▶ by sampling  $F$
- ▶ by using Rademacher complexity and Vapnik Vapnik-Chervonenkis theory.

This classical bound on the statistical error does not guarantee that the generalization error is small when  $\log(|F|)$  is large.

## Interpolation regime

- ▶ If the number of parameters is sufficient, it is possible to have  $R_n(f_n) = 0$ .
- ▶ In that case, it seems that the statistical error might to be way smaller than the previous bound.

For instance, for Wide Resnet,  $\frac{p}{n} = 179$  with

- ▶  $p$  : number of parameters in the network
- ▶  $n$  : number of samples

# Double descent

[Belkin et al., 2019]

## Optimization of neural networks

Difficulties of optimizing neural networks

Specific methods for neural networks

## Statistical learning

Bounding the estimation error

Interpolation regime and double descent

## Local averaging methods

Supervised learning

Density estimation

## Adaptivity

No free lunch theorems

Adaptivity



## Local averaging methods

Local averaging methods : approximation **without** optimization of an empirical risk.

## Setting

Let  $l$  be a loss. Generalization error :

$$R(f) = E_{(X,Y) \sim \rho}[l(Y, f(X))] \quad (28)$$

Bayes estimator :

$$f^* = \arg \min_{f \text{ measurable}} R(f) \quad (29)$$

Bayes risk :

$$R^* = E_X \left[ \inf_{y \in \mathcal{Y}} E_{Y \sim dP(Y|X)} [l(Y, y) | X] \right] \quad (30)$$

As always, the law of  $(X, Y)$  and of  $Y|X$  are unknown.

## Local averaging

### Local averaging :

- ▶ based on the dataset  $D_n$ , compute an approximation of the law  $Y|X$  or  $(Y, X)$ , without optimization of an empirical risk.
- ▶ use this approximation in the estimator (e.g. to compute an approximate conditional expected value.)

## Local averaging : regression

$\tilde{f}(x)$  : local averaging estimator, in the case of regression, squared loss, we can use

$$\tilde{f}(x) = \int_{y \in \mathbb{R}} y d\hat{P}(Y|X = x) \quad (31)$$

## Local averaging : classification

$\tilde{f}(x)$  : local averaging estimator, in the case of binary classification, squared loss, we can use

$$\tilde{f}(x) = \arg \max_{z \in \mathcal{Y}} \hat{P}(Y = z | X = x) \quad (32)$$

## Linear estimators

The question is then : how to choose the approximation  $\hat{d}P(Y|X=x)$ ?

**Linear estimators**

$$\hat{d}P(Y|X=x) = \sum_{i=1}^n \hat{w}_i(x) \delta_{y_i} \quad (33)$$

$\delta_{y_i}$  is the Dirac mass in  $y_i$ .

- ▶  $\forall i, \hat{w}_i(x) \geq 0$
- ▶  $\sum_{i=1}^n \hat{w}_i(x) = 1$

## Linear estimators

### Linear estimators

$$\hat{dP}(Y|X=x) = \sum_{i=1}^n \hat{w}_i(x) \delta_{y_i} \quad (34)$$

$\delta_{y_i}$  is the Dirac mass in  $y_i$ .

- ▶  $\forall i, \hat{w}_i(x) \geq 0$
- ▶  $\sum_{i=1}^n \hat{w}_i(x) = 1$

Application to regression :

$$\tilde{f}(x) = \sum_{i=1}^n \hat{w}_i(x) y_i \quad (35)$$

## Linear estimators

### Linear estimators

$$\hat{dP}(Y|X=x) = \sum_{i=1}^n \hat{w}_i(x) \delta_{y_i} \quad (36)$$

$\delta_{y_i}$  is the Dirac mass in  $y_i$ .

- ▶  $\forall i, \hat{w}_i(x) \geq 0$
- ▶  $\sum_{i=1}^n \hat{w}_i(x) = 1$

Application to classification :

$$\tilde{f}(x) = \arg \max_{j \in \{0,1\}} \sum_{i=1}^n \hat{w}_i(x) 1_{y_i=j} \quad (37)$$



## Choice of the weights

### Linear estimators

$$\hat{dP}(Y|X=x) = \sum_{i=1}^n \hat{w}_i(x) \delta_{y_i}(y) \quad (38)$$

$\delta_{y_i}$  is the Dirac mass in  $y_i$ .

- ▶  $\forall i, \hat{w}_i(x) \geq 0$
- ▶  $\sum_{i=1}^n \hat{w}_i(x) = 1$

For any sample  $i$ , the weight function  $\hat{w}_i(x)$  should be

- ▶ closer to 1 for training point  $x_i$  that are close to  $x$ .
- ▶ closer to 0 for training point  $x_i$  that are far from  $x$ .

# Choice of the weights

## Linear estimators

$$\hat{d}P(Y|X=x) = \sum_{i=1}^n \hat{w}_i(x) \delta_{y_i}(y) \quad (39)$$

$\delta_{y_i}$  is the Dirac mass in  $y_i$ .

- ▶  $\forall i, \hat{w}_i(x) \geq 0$
- ▶  $\sum_{i=1}^n \hat{w}_i(x) = 1$


Three possibilities :

- ▶ partition estimators
- ▶ nearest neighbors
- ▶ Nadaraya-Watson (kernel regression)


## Nearest neighbors

Given  $k \geq 1$ , and a metric  $d$  on  $\mathcal{X}$ , average the predictions of the  $k$  nearest neighbors (for regression) or take the majority vote (for classification).


## Nearest neighbors




## Nearest neighbors




## Nearest neighbors



## Nearest neighbors




# Nearest neighbors







## Nearest neighbors



# Nearest neighbors



## Nearest neighbors



## Nearest neighbors

Given  $k \geq 1$ , and a metric  $d$  on  $\mathcal{X}$ , average the predictions of the  $k$  nearest neighbors (for regression) or take the majority vote (for classification).

**Exercise 5 :** What is  $\hat{w}_i(x)$  ?

## Nearest neighbors

$$\hat{w}_i(x) : \begin{cases} 1/k & \text{if } i \text{ is in the closest neighbors} \\ 0 & \text{otherwise} \end{cases}$$

## Nearest neighbors

$k$  is a hyperparameter, hence it must be tuned, for instance with cross validation.

- ▶ too small  $k$  : underfitting
- ▶ too large  $k$  : overfitting

## Nearest neighbors search

The search for nearest neighbors is a problem itself!

<https://scikit-learn.org/stable/modules/neighbors.html>

[https://en.wikipedia.org/wiki/K-d\\_tree](https://en.wikipedia.org/wiki/K-d_tree)

[https://en.wikipedia.org/wiki/Ball\\_tree](https://en.wikipedia.org/wiki/Ball_tree)

## Partition estimators

$$\mathcal{X} = \cup_{j \in J} A_j.$$

$A_1$	$A_2$	$A_3$	$A_4$	$A_5$
$A_6$	$A_7$	$A_8$	$A_9$	$A_{10}$
$A_{11}$	$A_{12}$	$A_{13}$	$A_{14}$	$A_{15}$
$A_{16}$	$A_{17}$	$A_{18}$	$A_{19}$	$A_{20}$
$A_{21}$	$A_{22}$	$A_{23}$	$A_{24}$	$A_{25}$

For each  $x$ , average the predictions of the samples that are in the same  $A_j$  as  $x$ . We can note it  $A(x)$ .



## Partition estimators

$$\mathcal{X} = \cup_{j \in J} A_j.$$

$A_1$	$A_2$	$A_3$	$A_4$	$A_5$
$A_6$	$A_7$	$A_8$	$A_9$	$A_{10}$
$A_{11}$	$A_{12}$	$A_{13}$	$A_{14}$	$A_{15}$
$A_{16}$	$A_{17}$	$A_{18}$	$A_{19}$	$A_{20}$
$A_{21}$	$A_{22}$	$A_{23}$	$A_{24}$	$A_{25}$

For each  $x$ , average the predictions of the samples that are in the same  $A_j$  as  $x$ . We can note it  $A(x)$ .

**Exercice 6 :** What is  $\hat{w}_i(x)$  ?

## Partition estimator

$$\hat{w}_i(x) = \frac{1_{x_i \in A(x)}}{\sum_{k=1}^n 1_{x_k \in A(x)}} \quad (40)$$

## Partition estimator

**Exercise 7 :** We have seen in previous classes one example of partition estimator. What is it ?

## Kernel regression (Nadaraya-Watson)

We consider a non-negative kernel function  $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}_+$  and

$$\hat{w}_i(x) = \frac{k(x, x_i)}{\sum_{i=1}^n k(x, x_i)} \quad (41)$$

## Non-negative kernels

Often

$$k(x, x') = \frac{1}{h^d} q\left(\frac{x - x'}{h}\right) \quad (42)$$

with  $d$  the dimension,  $h$  a bandwidth parameter.




Image from [Bach, 2021].

## Non-negative kernels

$$k(x, x') = \frac{1}{h^d} q\left(\frac{x - x'}{h}\right) \quad (43)$$




Image from [Bach, 2021].

- ▶ Box kernel :  $q(x) = 1_{||x|| \leq 1}$
- ▶ Gaussian kernel :  $q(x) = e^{-\frac{||x||^2}{2}}$

## Remark

These kernels are not exactly the same as the ones we mentioned earlier (positive-definite kernels).

These kernels are more simply non-negative (less specific).

Estimator :

$$f(\tilde{x}) = \frac{\sum_{i=1}^n k(x, x_i) y_i}{\sum_{i=1}^n k(x, x_i)} \quad (44)$$

## Curse of dimensionality

It is possible to show, that under some simple regularity assumptions on the target, the convergence rate of the error of these estimators, as a function of  $n$ , is  $\mathcal{O}(n^{-\frac{2}{d+2}})$ , where  $d$  is the underlying dimension.

In order to have an error smaller than  $\epsilon$ , we need to have

$$n \geq \left(\frac{1}{\epsilon}\right)^{\frac{d+2}{2}} \quad (45)$$

- ▶ It is not easy to exploit a higher regularity of the target function ( **no adaptivity to the regularity** )
- ▶ It is not possible to learn with these methods in high dimension.



## Kernel density estimation

It is possible to use similar ideas to perform Kernel density estimation (KDE). (Again, here it is just a non-negative kernel)

<https://francisbach.com/cursed-kernels/>

[https:](https://seaborn.pydata.org/generated/seaborn.jointplot.html)

[//seaborn.pydata.org/generated/seaborn.jointplot.html](https://seaborn.pydata.org/generated/seaborn.jointplot.html)

[https://fr.wikipedia.org/wiki/Estimation\\_par\\_noyau](https://fr.wikipedia.org/wiki/Estimation_par_noyau)

[https:](https://en.wikipedia.org/wiki/Kernel_density_estimation)

[//en.wikipedia.org/wiki/Kernel\\_density\\_estimation](https://en.wikipedia.org/wiki/Kernel_density_estimation)

## Optimization of neural networks

- Difficulties of optimizing neural networks

- Specific methods for neural networks

## Statistical learning

- Bounding the estimation error

- Interpolation regime and double descent

## Local averaging methods

- Supervised learning

- Density estimation

## Adaptivity

- No free lunch theorems

- Adaptivity

## No free lunch theorems

$\mathcal{A}$  : learning rule. Takes the dataset  $D_n$  as input and outputs an estimator  $f_n$  (for instance based on empirical risk minimization, local averaging, etc).

There are several no free lunch theorems.

# No free lunch theorems

## Theorem

*No free lunch - fixed  $n$*

*We consider a binary classification task with "0-1"-loss, and  $\mathcal{X}$  infinite.*

*We note  $\mathcal{P}$  the set of all probability distributions on  $\mathcal{X} \times \{0, 1\}$ .*

*For any  $n > 0$  and any learning rule  $\mathcal{A}$*

$$\sup_{dp \in \mathcal{P}} E \left[ R_{dp}(\mathcal{A}(D_n(dp))) \right] - R_{dp}^* \geq \frac{1}{2} \quad (46)$$

We write  $D_n(dp)$  in order to emphasize that the dataset is sampled randomly from the distribution  $dp$ .

## No free lunch

- ▶ For any learning rule, there exists a distribution for which this learning rule performs badly.
- ▶ No method is universal and can have a good convergence rate on all problems.

**However**, considering **all** problems is probably not relevant for machine learning.

# Adaptivity

If the learning rule improves (faster convergence rate) when we add a property on the problem (for instance, regularity of the target function), we say that we have adaptivity to this property.

For instance : gradient descent is adaptive to the strong convexity of the target function, since with a proper choice of the learning rate  $\gamma$ , the convergence rate is exponential, with a rate that involves the strong convexity constant  $\mu$ .

There are several forms of adaptivity.

## Most general case

The target is just Lipschitz-continuous, no extra-hypothesis. In this case the optimal rate is of the form  $\mathcal{O}(n^{-\frac{1}{d}})$  (curse of dimensionality) for all learning rules.

## Adaptivity to the input space

If the input data lie on a submanifold (e.g. a subspace) of  $\mathbb{R}^d$  of lower dimension than  $d$ , most methods adapt to this property.



## Adaptivity to the regularity of the target function

If the target is smoother (meaning that all derivatives up to order  $m$  are bounded), kernel methods (here, positive-definite kernels) and neural network adapt, if well optimized and regularized. The rate can become  $\mathcal{O}(n^{-\frac{m}{d}})$ .

## Adaptivity to latent variables

If the target function depends only on a  $k$  dimensional linear projection of the data, neural networks adapt, if well optimized. The rate can become  $O(n^{-\frac{m}{k}})$ .

<https://francisbach.com/quest-for-adaptivity/>

# References I



Bach, F. (2021).

Learning Theory from First Principles Draft.

*Book Draft*, page 229.



Belkin, M., Hsu, D., Ma, S., and Mandal, S. (2019).

Reconciling modern machine-learning practice and the classical bias–variance trade-off.

*Proceedings of the National Academy of Sciences of the United States of America*, 116(32) :15849–15854.